

Assignment 2: Web scrapping/Data Wrangling in R

Suman Paudel 33

2024-03-09

GitHub Repository

Code Execution and Output/Interpretation of Session 7

What is a “package” in R?

- In R, the fundamental unit of shareable code/function is the package.
- A package bundles together code, data, documentation, and tests, and is easy to share with others.
- The public clearing house for R packages is the Comprehensive R Archive Network or CRAN. The details about the packages can be found at CRAN.

How to install, use and get help about any package from CRAN?

- We can install packages of CRAN in R using: `install.packages("packagename")`
For example: `install.packages("dplyr")`
- We can then use the installed packages in R using:
`library(packagename)`
e.g. `library(dplyr)`
- We can get help on the installed packages in R using:
`?packagename` or `help(package = "packagename")`
e.g. `?dplyr` OR `help(package = "dplyr")`

How to develop a package in R? <https://hilaryparker.com/2014/04/29/writing-an-r-package-from-scratch/>

- Step 0: Packages you will need
- Step 1: Creating your package directory
- Step 2: Add functions
- Step 3: Add documentation
- Step 4: Process your documentation
- Step 5: Install
- Step 6: Make a package GitHub repo (Bonus!)
- Step 7: Infinity- Iterate

Reading (Import) data in R/R Studio:

- Text files: R base, readr etc
- Excel files: readXL, openxls etc.
- SPSS, Stata, SAS files: foreign, haven etc.

Reading data in R/R Studio

- JSON files: rjson, jsonlite, RJSONIO etc.

To install the library to work on json files.

- `install.packages('jsonlite')`
- `install.packages('rjson')`

Code Sample 1

```
# vector
library("rjson")
data <- fromJSON(file = "json_data.json")
data

## $ID
## [1] "1" "2" "3" "4" "5" "6" "7" "8"
##
## $Name
## [1] "Rick"      "Dan"      "Michelle" "Ryan"      "Gary"      "Nina"      "Simon"
## [8] "Guru"
##
## $Salary
## [1] "623.3"  "515.2"  "611"     "729"     "843.25"  "578"     "632.8"  "722.5"
##
## $StartDate
## [1] "1/1/2012"  "9/23/2013" "11/15/2014" "5/11/2014" "3/27/2015"
## [6] "5/21/2013" "7/30/2013" "6/17/2014"
##
## $Dept
## [1] "IT"      "Operations" "IT"      "HR"      "Finance"
## [6] "IT"      "Operations" "Finance"
```

Interpretation:

- loaded the `json_data.json` file from current working directory using `rjson` library.

Code Sample 2

```
# covert to data frame:
json_data_frame <- as.data.frame(data)
json_data_frame

##   ID    Name Salary StartDate    Dept
## 1  1    Rick  623.3   1/1/2012      IT
## 2  2     Dan  515.2   9/23/2013 Operations
## 3  3 Michelle   611 11/15/2014      IT
## 4  4     Ryan   729   5/11/2014      HR
## 5  5     Gary 843.25  3/27/2015  Finance
## 6  6     Nina   578   5/21/2013      IT
## 7  7    Simon 632.8  7/30/2013 Operations
## 8  8     Guru 722.5   6/17/2014  Finance
```

Interpretation:

- converted the loaded json data into data.frame

Code Sample 3

```
# summary of jason_data_frame
summary(jason_data_frame)
```


```
##      ID           Name           Salary           StartDate
## Length:8       Length:8       Length:8       Length:8
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##      Dept
## Length:8
## Class :character
## Mode  :character
```

Interpretation:

- summarizes the json data frame.

Code Sample

```
# histogram of salary of jason data frame
# since salary was of character so hist was not produced, so I converted it to integer.
jason_data_frame$Salary <- as.integer(jason_data_frame$Salary)
hist(jason_data_frame$Salary)
```




Interpretation:

- since salary was of character so hist was not produced, so I converted it to integer.
- generates histogram of salary of jason data frame

Code Sample

```
# histogram of salary of jason data frame  
# since salary was of character so hist was not produced, so I converted it to integer.  
jason_data_frame$Salary <- as.integer(jason_data_frame$Salary)  
hist(jason_data_frame$Salary)
```

Histogram of jason_data_frame\$Salary



Interpretation:

- since salary was of character so hist was not produced, so I converted it to integer.
- generates histogram of salary of jason data frame

Code Sample

```
# average salary by department
aggregate(Salary ~ Dept, data = jason_data_frame, FUN = mean)
```

```
##      Dept Salary
## 1  Finance  782.5
## 2      HR   729.0
## 3      IT   604.0
## 4 Operations  573.5
```

Interpretation:

- aggregates the average salary by department.

Code Sample

```
# also using dplyr package
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(magrittr)
avg_by_dept <-
  jason_data_frame %>%
  group_by(Dept) %>%
  summarise(avg_salart = mean(Salary))
```

Interpretation:

- average salary by department can be achieved by using dplyr package and using summarise.

Code Sample

```
# frequency distribution of all variables
# since tables only shows the frequency distribution of factors, table on numerical or integer value is
table(jason_data_frame$Dept)

##
##   Finance      HR      IT Operations
##         2         1         3         2
```

Interpretation:

- since tables only shows the frequency distribution of factors, table on numerical or integer value is bad practice.

Reading JSON file from URL: Web API

Code Sample

```
# fetch the data from web api
library(jsonlite)
raw <-
  fromJSON(
    "https://data.ny.gov/api/views/9a8c-vfzj/rows.json?accessType=DOWNLOAD")
food_market <- raw[['data']]
```

Interpretation:

- read the json data from API using jsonlite library
- the raw object gets the list of values from where data part is extracted and assigned to food_market object

Code Sample

```
# structure and head of food market
str(food_market)

## chr [1:24196, 1:24] "row-4wb8.javd.fc2j" "row-cxk5.daxz~ky99" ...

head(food_market, n = 4)

##      [,1]                [,2]                [,3]
## [1,] "row-4wb8.javd.fc2j" "00000000-0000-0000-5489-E0BA9784EE18" "0"
## [2,] "row-cxk5.daxz~ky99" "00000000-0000-0000-1138-77B40E28ADAD" "0"
## [3,] "row-mgqk~8t89~nkze" "00000000-0000-0000-29DB-90F34B446A8B" "0"
## [4,] "row-98en.nphr~vg8m" "00000000-0000-0000-B908-B0E6F9A8F452" "0"
##      [,4]      [,5] [,6]      [,7] [,8]      [,9]      [,10]      [,11]      [,12]
## [1,] "1707244443" NA   "1707244449" NA   "{ }" "BRONX" "750994" "Store" "JAC"
## [2,] "1707244443" NA   "1707244449" NA   "{ }" "BRONX" "608900" "Store" "JAC"
## [3,] "1707244443" NA   "1707244449" NA   "{ }" "BRONX" "607918" "Store" "JAC"
## [4,] "1707244443" NA   "1707244449" NA   "{ }" "BRONX" "746827" "Store" "JAC"
##      [,13]      [,14]      [,15] [,16]      [,17] [,18] [,19]
## [1,] "CA FOOD MART INC" "CA FOOD MART" "3706" "3RD AVENUE" NA   NA   "BRONX"
## [2,] "1307 CORP"       "1307 CORP"     "1307" "CROSBY AVE" NA   NA   "BRONX"
## [3,] "GNP SUNIL CORP"  "GNP SUNIL"     "585"  "W 235TH ST" NA   NA   "BRONX"
## [4,] "BRONX BAZAR INC" "BRONX BAZAR"   "2550" "BOSTON RD"  NA   NA   "BRONX"
##      [,20] [,21]      [,22]      [,23]                [,24]
## [1,] "NY"   "10456" "0"     "POINT (-73.903143068 40.835323772)" "307"
## [2,] "NY"   "10461" "600"    "POINT (-73.82918325 40.838483414)" "307"
## [3,] "NY"   "10463" "1500"   "POINT (-73.911162035 40.885983701)" "307"
## [4,] "NY"   "10467" "0"     "POINT (-73.863887262 40.8642032)"  "307"
```

Interpretation:

- returns the structure and head of food_market.

Code Sample

```
names <- food_market[,14]
```

Interpretation:

- returns the names stores listed in index 14 of food_market in names object

Code Sample

```
head(names)

## [1] "CA FOOD MART" "1307 CORP"      "GNP SUNIL"      "BRONX BAZAR"  "TOPS 754"
## [6] "MAY STORE"

head(table(names), n = 7)
```

```
## names
## #248-PRICE RITE OF SOUTH AVE      #249 PRICE RITE OF UTICA
##                                1                                1
##                                @JENSARTISAN                      0113 FOOD
##                                1                                1
##                                056 DELI & GROCERY                1-800-FLOWERS
##                                1                                1
##                                1 BASKET MARKET
##                                1
```

Interpretation:

- `head(names)` returns first few data.
- `head(table(names), n = 7)` returns the frequency table of names of stores.
- since data is so much I restricted to first 7 values for frequency table.

Code Sample

```
table(v19)

## Error in eval(expr, envir, enclos): object 'v19' not found
```

Interpretation:

- trying to get the frequency table v19. but v19 doesn't exist and throws error

Code Sample

```
table(food_market$V19)

## Error in food_market$V19: $ operator is invalid for atomic vectors
```

Interpretation:

- An atomic vector is a one-dimensional data object created using functions like `c()` or `vector()` in R.
- The `$` operator is commonly used to access elements of a data frame by column name.
- since `food_market` is of atomic vectors type, while access the element using `$` returned this error.
- this can be mitigated by using `table(food_market[,19])`

Code Sample

```
head(table(food_market[, 19]), n = 5)

##
## ACCORD      ADAMS ADAMS CENTER      ADDISON      ADIRONDACK
##          5          9          2          8          1
```

Interpretation:

- returns the frequency table of data that is at 19th column of `food_market`
- since data is too much, I restricted the display to 5.

Web scrapping in R: A Simple (barebones) Example

Code Sample

```
# The recommended package for web scrapping in R is "rvest"
# using html
library(rvest)
simple <-
  read_html('https://dataquestio.github.io/web-scraping-pages/simple.html')
simple %>% html_nodes('p') %>% html_text()

## [1] "Here is some simple content for this page."
```

Interpretation:

- reads the html using `read_html()`
- we loaded the html in simple object then using pipes parsed the html with node 'p' then converted that to normal text.

Code Sample

```
# Wikipedia Nepal Covid 19
wiki_link <- "https://en.wikipedia.org/wiki/COVID-19_pandemic_in_Nepal"
wiki_page <- read_html(wiki_link)
```

Interpretation:

- since salary was of character so hist was not produced, so I converted it to integer.
- generates histogram of salary of jason data frame

Code Sample

```
str(wiki_page)

## List of 2
## $ node:<externalptr>
## $ doc :<externalptr>
## - attr(*, "class")= chr [1:2] "xml_document" "xml_node"
```

Interpretation:

- structure of wiki page, returns list of 2 xml node and xml doc

Code Sample

```
wiki_page %>% html_nodes("table")

## {xml_nodeset (68)}
## [1] <table class="box-Update plainlinks metadata ambox ambox-content ambox-U ...
## [2] <table class="infobox"><tbody>\n<tr><th colspan="2" class="infobox-above ...
## [3] <table>\n<caption class="bb-default">\nCOVID-19 cases in Nepal<span clas ...
## [4] <table class="wikitable floatright mw-collapsible {{{class}}}" style="wi ...
## [5] <table class="box-Update plainlinks metadata ambox ambox-content ambox-U ...
## [6] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [7] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [8] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [9] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [10] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [11] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [12] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [13] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [14] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [15] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [16] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [17] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [18] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [19] <table class="box-Notice plainlinks metadata ambox ambox-notice" role="p ...
## [20] <table class="box-Update plainlinks metadata ambox ambox-content ambox-U ...
## ...
```

Interpretation:

- parses the html nodes containing `table` in html page of `wiki_page`.

Code Sample

```
covid_table <-
  wiki_page %>% html_elements('.COVID-19_pandemic_data_Nepal_medical_cases') %>% html_node('table') %>%
covid_table <- covid_table[[1]]
```

Interpretation:

- from wiki page, first parsed the html elements having `.COVID-19_pandemic_data_Nepal_medical_cases`
- then extracted the node containing `table` into `html_table()`
- since the data was returned in the list, the actual data was at first index which is dataframe.

Data wrangling: Part I Column names of covid_table

Code Sample

```
names(covid_table) <-
  paste(names(covid_table), covid_table[1, ], sep = "_")
covid_table <- covid_table[-1, ]
names(covid_table)
```

```
## [1] "Date_Date"          "Confirmed cases_Total" "Confirmed cases_New"
## [4] "Confirmed cases_Active" "Recoveries_Total"     "Recoveries_New"
## [7] "Deaths_Total"        "Deaths_New"           "RT-PCR tests_Total"
## [10] "RT-PCR tests_New"    "TPR_TPR"              "RR_RR"
## [13] "CFR_CFR"            "Ref._Ref."            "
```

Interpretation:

- even though the data was extracted to dataframe there was error in columns and we want to rename it
- concatenated the column names with first row of dataframe.
- and finally dropping the first column.

Code Sample

```
#Check the structure of data again
str(covid_table)

## tibble [495 x 14] (S3: tbl_df/tbl/data.frame)
##  $ Date_Date          : chr [1:495] "23 Jan" "24 Jan" "25 Jan" "26 Jan" ...
##  $ Confirmed cases_Total : chr [1:495] "1" "1" "1" "1" ...
##  $ Confirmed cases_New   : chr [1:495] "+1" "0" "0" "0" ...
##  $ Confirmed cases_Active: chr [1:495] "1" "1" "1" "1" ...
##  $ Recoveries_Total      : chr [1:495] "0" "0" "0" "0" ...
##  $ Recoveries_New       : chr [1:495] "0" "0" "0" "0" ...
##  $ Deaths_Total        : chr [1:495] "0" "0" "0" "0" ...
##  $ Deaths_New          : chr [1:495] "0" "0" "0" "0" ...
##  $ RT-PCR tests_Total    : chr [1:495] "" "" "" "" ...
##  $ RT-PCR tests_New     : chr [1:495] "" "" "" "" ...
##  $ TPR_TPR              : chr [1:495] "" "" "" "" ...
##  $ RR_RR                : chr [1:495] "0%" "0%" "0%" "0%" ...
##  $ CFR_CFR              : chr [1:495] "0%" "0%" "0%" "0%" ...
##  $ Ref._Ref.            : chr [1:495] "[175]" "" "" "" ...
```

Interpretation:

- checking the structure again to ensure.

Data wrangling: Part II Renaming the column names

- Change “Date_Date” variable as “**Date**”
- Change “Confirmed cases_Total” variable as “**Confirmed_Cases_Total**”
- Change “Confirmed cases_New” variable as “**Confirmed_Cases_New**”
- Change “Confirmed cases_Active” variable as “**Confirmed_Cases_Active**”
- Change “RT-PCR tests_Total” variable as “**RT-PCR_tests_Total**”
- Change “RT-PCR tests_New” variable as “**RT-PCR_tests_New**”
- Change “TPR_TPR” variable as “**TPR**”
- Change “RR_RR” variable as “**RR**”
- Change “CFR_CFR” variable as “**CFR**”
- Change “Ref._Ref.” variable as “**Ref**”

Code Sample

```
colnames(covid_table) <- c(
  "Date",
  "Confirmed_Cases_Total",
  "Confirmed_Cases_New",
  "Confirmed_Cases_Active",
  "Recoveries_Total",
  "Recoveries_New",
  "Deaths_Total",
  "Deaths_New",
  "PCR_Total",
  "PCR_New",
  "TPR",
  "RR",
  "CFR",
  "Ref"
)
```

Interpretation:

- renaming the column names to desired column names.

Code Sample

```
colnames(covid_table) <- c(
  "Date",
  "Confirmed_Cases_Total",
  "Confirmed_Cases_New",
  "Confirmed_Cases_Active",
  "Recoveries_Total",
  "Recoveries_New",
  "Deaths_Total",
  "Deaths_New",
  "PCR_Total",
  "PCR_New",
  "TPR",
  "RR",
  "CFR",
  "Ref"
)
```

Interpretation:

- renaming the column names to desired column names.

Code Sample

```
# another way to do same thing

names(covid_table)[names(covid_table)=='Date_Date'] = 'Date'
names(covid_table)[names(covid_table)=='Confirmed cases_Total'] = 'Confirmed_Cases_Total'
```

```
names(covid_table)[names(covid_table)=='Confirmed cases_New'] = 'Confirmed_Cases_New'
names(covid_table)[names(covid_table)=='Confirmed cases_Active'] = 'Confirmed_Cases_Active'
names(covid_table)[names(covid_table)=='Recoveries_Total'] = 'Recoveries_Total'
names(covid_table)[names(covid_table)=='Recoveries_New'] = 'Recoveries_New'
names(covid_table)[names(covid_table)=='Deaths_Total'] = 'Deaths_Total'
names(covid_table)[names(covid_table)=='Deaths_New'] = 'Deaths_New'
names(covid_table)[names(covid_table)=='RT-PCR tests_Total'] = 'PCR_Total'
names(covid_table)[names(covid_table)=='RT-PCR tests_New'] = 'PCR_New'
names(covid_table)[names(covid_table)=='TPR_TPR'] = 'TPR'
names(covid_table)[names(covid_table)=='RR_RR'] = 'RR'
names(covid_table)[names(covid_table)=='CFR_CFR'] = 'CFR'
names(covid_table)[names(covid_table)=='Ref._Ref.'] = 'Ref'
```

Interpretation:

- Another way to rename the column names to desired column names.

Data Wrangling Part III

Code Sample

```
# removing "+" and "%" from the variables
covid_table$Confirmed_Cases_New <- gsub('[+]', '', covid_table$Confirmed_Cases_New)
covid_table$Recoveries_New <- gsub('[+]', '', covid_table$Recoveries_New)
covid_table$Deaths_New <- gsub('[+]', '', covid_table$Deaths_New)
covid_table$PCR_New <- gsub('[+]', '', covid_table$PCR_New)

covid_table$TPR <- gsub('[%]', '', covid_table$TPR)
covid_table$RR <- gsub('[%]', '', covid_table$RR)
covid_table$CFR <- gsub('[%]', '', covid_table$CFR)
```

Interpretation:

- cleaning the messy data using gsub.
- removing % and + from data.

Code Sample

```
# since in pdf while converting chr to numeric or integer type coerce warning was given for invalid typ
covid_table$Confirmed_Cases_Total <- gsub('[,]', '', covid_table$Confirmed_Cases_Total)
covid_table$Confirmed_Cases_New <- gsub('[,]', '', covid_table$Confirmed_Cases_New)
covid_table$Confirmed_Cases_Active <- gsub('[,]', '', covid_table$Confirmed_Cases_Active)
covid_table$Recoveries_Total <- gsub('[,]', '', covid_table$Recoveries_Total)
covid_table$Recoveries_New <- gsub('[,]', '', covid_table$Recoveries_New)

covid_table$Deaths_Total <- gsub('[,]', '', covid_table$Deaths_Total)
covid_table$Deaths_New <- gsub('[,]', '', covid_table$Deaths_New)
covid_table$PCR_Total <- gsub('[,]', '', covid_table$PCR_Total)
covid_table$PCR_New <- gsub('[,]', '', covid_table$PCR_New)

covid_table$TPR <- gsub('[,]', '', covid_table$TPR)
covid_table$RR <- gsub('[,]', '', covid_table$RR)
covid_table$CFR <- gsub('[,]', '', covid_table$CFR)
```

Interpretation:

- Since while converting chr to integer type it gave some warning about coerce where values were invalid type for e.g 313,333 can't be interpreted as integer because of comma. But it is right in context of html in website.
- removed the , using gsub.

Data wrangling: Part IV Converting “chr” variables as numbers

Code Sample

```
covid_table$Confirmed_Cases_Total <- as.integer(covid_table$Confirmed_Cases_Total)
covid_table$Confirmed_Cases_New <- as.integer(covid_table$Confirmed_Cases_New)
covid_table$Confirmed_Cases_Active <- as.integer(covid_table$Confirmed_Cases_Active)
covid_table$Recoveries_Total <- as.integer(covid_table$Recoveries_Total)
covid_table$Recoveries_New <- as.integer(covid_table$Recoveries_New)

covid_table$Deaths_Total <- as.integer(covid_table$Deaths_Total)
covid_table$Deaths_New <- as.integer(covid_table$Deaths_New)
covid_table$PCR_Total <- as.integer(covid_table$PCR_Total)
covid_table$PCR_New <- as.integer(covid_table$PCR_New)

covid_table$TPR <- as.numeric(covid_table$TPR)
covid_table$RR <- as.numeric(covid_table$RR)
covid_table$CFR <- as.numeric(covid_table$CFR)
```

Interpretation:

- converted all of the required columns of chr type of integer or numeric type using `as.numeric()` function.

Code Sample

```
str(covid_table)

## tibble [495 x 14] (S3: tbl_df/tbl/data.frame)
##  $ Date                : chr [1:495] "23 Jan" "24 Jan" "25 Jan" "26 Jan" ...
##  $ Confirmed_Cases_Total : int [1:495] 1 1 1 1 1 1 1 1 1 1 ...
##  $ Confirmed_Cases_New   : int [1:495] 1 0 0 0 0 0 0 0 0 0 ...
##  $ Confirmed_Cases_Active: int [1:495] 1 1 1 1 1 1 0 0 0 0 ...
##  $ Recoveries_Total      : int [1:495] 0 0 0 0 0 0 1 1 1 1 ...
##  $ Recoveries_New       : int [1:495] 0 0 0 0 0 0 1 0 0 0 ...
##  $ Deaths_Total        : int [1:495] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Deaths_New          : int [1:495] 0 0 0 0 0 0 0 0 0 0 ...
##  $ PCR_Total            : int [1:495] NA NA NA NA NA 3 4 5 5 NA ...
##  $ PCR_New              : int [1:495] NA NA NA NA NA NA 1 1 0 NA ...
##  $ TPR                  : num [1:495] NA NA NA NA NA ...
##  $ RR                   : num [1:495] 0 0 0 0 0 0 100 100 100 100 ...
##  $ CFR                  : num [1:495] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Ref                  : chr [1:495] "[175]" "" "" "" "" ...
```

Interpretation:

- finally check the data type.

Converting Date (Assignment)

Code Sample

```
head(as.Date(covid_table$Date, format = "%d %b"), n = 50)

## [1] "2024-01-23" "2024-01-24" "2024-01-25" "2024-01-26" "2024-01-27"
## [6] "2024-01-28" "2024-01-29" "2024-01-30" "2024-01-31" "2024-02-01"
## [11] "2024-02-02" "2024-02-03" "2024-02-04" "2024-02-05" "2024-02-06"
## [16] "2024-02-07" "2024-02-08" "2024-02-09" "2024-02-10" "2024-02-11"
## [21] "2024-02-12" "2024-02-13" "2024-02-14" "2024-02-15" "2024-02-16"
## [26] "2024-02-17" "2024-02-18" "2024-02-19" "2024-02-20" "2024-02-21"
## [31] "2024-02-22" "2024-02-23" "2024-02-24" "2024-02-25" "2024-02-26"
## [36] "2024-02-27" "2024-02-28" "2024-02-29" "2024-03-01" "2024-03-02"
## [41] "2024-03-03" "2024-03-04" "2024-03-05" "2024-03-06" "2024-03-07"
## [46] "2024-03-08" "2024-03-09" "2024-03-10" "2024-03-11" "2024-03-12"
```

Interpretation:

- By default tries to check this format "%Y-%m-%d" or "%Y/%m/%d" for given date or string data.
- since our date column has day and month I have to use %d for day and %b for short name of month.
- returns the value in 2024 but that date is of 2020 and 2021
- In order to work around this I have made a logic below.

Code Sample

```
# using base R date time functions

start_date_2020 <- "2020-01-23"
end_date_2020 <- "2020-12-31"

# check days diff
days_diff <-
  as.numeric(difftime(end_date_2020, start_date_2020, units = 'days')) + 1

# check difference till first n rows and insert value if condition is true return 2020 else 2021
check_year <- ifelse(1:nrow(covid_table) <= days_diff, 2020, 2021)

# format the date according to need
formatted_date <- as.Date(covid_table$Date, format = "%d %b")
formatted_date <- format(formatted_date, '-%m-%d')

# concatenate the year and formatted day and date
final_converted_date <- paste0(check_year, formatted_date)

# finally since concatenation returns string again convert the type into date and load the Date column
covid_table$Date <- as.Date(final_converted_date)
head(covid_table$Date)

## [1] "2020-01-23" "2020-01-24" "2020-01-25" "2020-01-26" "2020-01-27"
## [6] "2020-01-28"
```

```
class(covid_table$Date)
```

```
## [1] "Date"
```

Interpretation:

- Since I know that data is of 23rd Jan 2020 to 31st May 2021.
- Also, after 31st Dec next year will arrive, so I set variable `start_date_2020 <- "2020-01-23"` and `end_date_2020 <- "2020-12-31"`.
- I calculated the difference in day from 23rd Jan to 31st Dec.
- Also data is serial in dataframe I filtered the data to 2020 and 2021 using difference in rows. Set 2020 to first 344 rows and rest 2021.
- since `as.Date(covid_table$Date, format = "%d %b")` gave the date of 2024, I formatted the date and only got day and month.
- Finally concatenated the calculated year with formatted day and month.

Code Sample

```
# using lubridate
library(lubridate)

start_date_2020 <- "2020-01-23"
end_date_2020 <- "2020-12-31"

# find difference
diff <-
  (ymd(start_date_2020) %--% ymd(end_date_2020)) %/% days(1) + 1

# check difference till first n rows and insert value if condition is true return 2020 else 2021
check_year <- ifelse(1:nrow(covid_table) <= diff, 2020, 2021)

# format the date to day and month
formatted_date <- ymd(as.Date(covid_table$Date, format = "%d %b"))

# create the date using make_date function
final_converted_date <-
  make_date(
    year = check_year,
    month = month(formatted_date),
    day = day(formatted_date)
  )

# change the Date column value with final converted date
covid_table$Date <- final_converted_date
class(covid_table$Date)

## [1] "Date"
```

Interpretation:

- Using lubridate to achieve the solution to given problem.

Code Execution and Output/Interpretation of Session 8 and 9

Code Sample

```
library(tidyverse)
library(magrittr)

table1 <- tibble(
  country = c("Afghanistan", "Afghanistan", "Brazil",
              "Brazil", "China", "China"),
  year = c(1999, 2000, 1999, 2000, 1999, 2000),
  cases = c(745, 2666, 37737, 80488, 212258, 213766),
  population = c(19987071, 20595360, 172006362,
                 174504898, 1272915272, 1280428583)
)
```

table1

```
## # A tibble: 6 x 4
##   country      year  cases population
##   <chr>      <dbl> <dbl>      <dbl>
## 1 Afghanistan 1999     745   19987071
## 2 Afghanistan 2000    2666   20595360
## 3 Brazil      1999   37737  172006362
## 4 Brazil      2000   80488  174504898
## 5 China       1999  212258 1272915272
## 6 China       2000  213766 1280428583
```

table2

```
## # A tibble: 12 x 4
##   country      year type      count
##   <chr>      <dbl> <chr>      <dbl>
## 1 Afghanistan 1999 cases         745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases         2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases         37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases         80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases         212258
## 10 China      1999 population 1272915272
## 11 China      2000 cases         213766
## 12 China      2000 population 1280428583
```

table3

```
## # A tibble: 6 x 3
##   country      year rate
##   <chr>      <dbl> <chr>
## 1 Afghanistan 1999 745/19987071
```

```
## 2 Afghanistan 2000 2666/20595360
## 3 Brazil      1999 37737/172006362
## 4 Brazil      2000 80488/174504898
## 5 China       1999 212258/1272915272
## 6 China       2000 213766/1280428583
```

table4a

```
## # A tibble: 3 x 3
##   country      '1999' '2000'
##   <chr>        <dbl> <dbl>
## 1 Afghanistan    745   2666
## 2 Brazil        37737  80488
## 3 China         212258 213766
```

table4b

```
## # A tibble: 3 x 3
##   country      '1999'      '2000'
##   <chr>        <dbl>      <dbl>
## 1 Afghanistan 19987071 20595360
## 2 Brazil      172006362 174504898
## 3 China       1272915272 1280428583
```

Interpretation:

- Exploring the few built in tables from tidyverse package.
- There are two main advantages of using tidy data:
 - There's a general advantage to picking one consistent way of storing data. If you have a consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
 - There's a specific advantage to placing variables in columns because it allows R's vectorized nature to shine.
 - dplyr, ggplot2, and all the other packages in the tidyverse are designed to work with tidy data.

Code Sample

```
# Pivoting - Longer to wider
# (To do standard statistical analysis)
table2 %>% pivot_wider(names_from = type, values_from = count)
```

```
## # A tibble: 6 x 4
##   country      year cases population
##   <chr>        <dbl> <dbl>      <dbl>
## 1 Afghanistan 1999    745   19987071
## 2 Afghanistan 2000   2666   20595360
## 3 Brazil      1999  37737  172006362
## 4 Brazil      2000  80488  174504898
## 5 China       1999 212258 1272915272
## 6 China       2000 213766 1280428583
```

Interpretation:

- Pivoting the table2 from longer to wider.
- usually used for standard statistical analysis

Code Sample

```
# Pivoting - Wider to Longer
# (To do Variance components analysis)
table4a %>%
  pivot_longer(c(`1999`, `2000`), names_to =
               "year", values_to = "cases")

## # A tibble: 6 x 3
##   country    year  cases
##   <chr>      <chr> <dbl>
## 1 Afghanistan 1999     745
## 2 Afghanistan 2000    2666
## 3 Brazil      1999   37737
## 4 Brazil      2000  80488
## 5 China       1999 212258
## 6 China       2000 213766
```

Interpretation:

- Pivoting the `table2` from wider to longer.
- usually used for standard statistical analysis.

Code Sample

```
# Tidy data: Separate
table3 %>% separate(rate, into = c("cases", "population"))

## # A tibble: 6 x 4
##   country    year cases population
##   <chr>      <dbl> <chr>   <chr>
## 1 Afghanistan 1999  745   19987071
## 2 Afghanistan 2000 2666   20595360
## 3 Brazil      1999 37737  172006362
## 4 Brazil      2000 80488  174504898
## 5 China       1999 212258 1272915272
## 6 China       2000 213766 1280428583

# OR
table3 %>% separate(rate, into = c("cases", "population"), sep = "/")

## # A tibble: 6 x 4
##   country    year cases population
##   <chr>      <dbl> <chr>   <chr>
## 1 Afghanistan 1999  745   19987071
## 2 Afghanistan 2000 2666   20595360
## 3 Brazil      1999 37737  172006362
## 4 Brazil      2000 80488  174504898
## 5 China       1999 212258 1272915272
## 6 China       2000 213766 1280428583
```

Interpretation:

- separate the value from single column to multiple columns.
- often useful when dealing with delimited values in single column.

Code Sample

```
# Tidy data: Unite
table5 %>% unite(new, century, year)

## # A tibble: 6 x 3
##   country    new    rate
##   <chr>      <chr> <chr>
## 1 Afghanistan 19_99 745/19987071
## 2 Afghanistan 20_00 2666/20595360
## 3 Brazil      19_99 37737/172006362
## 4 Brazil      20_00 80488/174504898
## 5 China       19_99 212258/1272915272
## 6 China       20_00 213766/1280428583

# OR
table5 %>% unite(new, century, year, sep = "")

## # A tibble: 6 x 3
##   country    new    rate
##   <chr>      <chr> <chr>
## 1 Afghanistan 1999  745/19987071
## 2 Afghanistan 2000  2666/20595360
## 3 Brazil      1999  37737/172006362
## 4 Brazil      2000  80488/174504898
## 5 China       1999  212258/1272915272
## 6 China       2000  213766/1280428583
```

Interpretation:

- unites columns into one.
- opposite of separate: it combines multiple columns into a single column.

Code Sample

```
# Missing values: Example
stocks <- tibble(
  year = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),
  qtr = c(1, 2, 3, 4, 2, 3, 4),
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)
)

stocks %>% pivot_wider(names_from = year, values_from = return)

## # A tibble: 4 x 3
##   qtr '2015' '2016'
##   <dbl> <dbl> <dbl>
## 1     1  1.88  NA
## 2     2  0.59  0.92
## 3     3  0.35  0.17
## 4     4  NA    2.66
```

Interpretation:

- pivoted the stocks to wider to check missing values on year and return
- The return for the fourth quarter of 2015 is explicitly missing, because the cell where its value should be instead contains NA.
- The return for the first quarter of 2016 is implicitly missing, because it simply does not appear in the dataset.

Code Sample

```
stocks %>%  
  pivot_wider(names_from = year, values_from = return) %>%  
  pivot_longer(  
    cols = c(`2015`, `2016`),  
    names_to = "year",  
    values_to = "return",  
    values_drop_na = TRUE  
  )
```

```
## # A tibble: 6 x 3  
##   qtr year  return  
##   <dbl> <chr>  <dbl>  
## 1     1 2015    1.88  
## 2     2 2015    0.59  
## 3     2 2016    0.92  
## 4     3 2015    0.35  
## 5     3 2016    0.17  
## 6     4 2016    2.66
```

Interpretation:

- `pivot_wider(names_from = year, values_from = return)` pivots data from longer to wider format. the date is spread out based on different years, where year is the column providing the year information, and return is the column providing return values.
- `pivot_longer()` is used to pivot data from wider to longer format. It gathers columns specified in cols into key-value pairs. Here, it specifies to gather columns 2015 and 2016 into a new column named “year”, with their corresponding values into a new column named “return”. The parameter `values_drop_na = TRUE` indicates that rows with missing values after pivoting should be dropped.

Code Sample

```
stocks %>% complete(year, qtr)
```

```
## # A tibble: 8 x 3  
##   year  qtr return  
##   <dbl> <dbl>  <dbl>  
## 1 2015     1   1.88  
## 2 2015     2   0.59  
## 3 2015     3   0.35  
## 4 2015     4    NA  
## 5 2016     1    NA
```

```
## 6 2016      2  0.92
## 7 2016      3  0.17
## 8 2016      4  2.66
```

Interpretation:

- `complete()` is used to complete a data frame with missing combinations of values. It takes columns specified in the arguments and ensures that all combinations of unique values are present in the resulting data frame.
- here, it ensures that all combinations of values from the columns `year` and `qtr` are present in the data frame, filling in any missing combinations with `NA` values.

Code Sample

```
treatment <- tribble(
  ~ person, ~ treatment, ~response,
  "Derrick Whitmore", 1, 7,
  NA, 2, 10,
  NA, 3, 9,
  "Katherine Burke", 1, 4
)
treatment
```

```
## # A tibble: 4 x 3
##   person      treatment response
##   <chr>          <dbl>     <dbl>
## 1 Derrick Whitmore      1         7
## 2 <NA>                 2        10
## 3 <NA>                 3         9
## 4 Katherine Burke      1         4
```

Interpretation:

- Create tibbles using an easier to read row-by-row layout. This is useful for small tables of data where readability is important.

Code Sample

```
treatment %>% fill(person)
```

```
## # A tibble: 4 x 3
##   person      treatment response
##   <chr>          <dbl>     <dbl>
## 1 Derrick Whitmore      1         7
## 2 Derrick Whitmore      2        10
## 3 Derrick Whitmore      3         9
## 4 Katherine Burke      1         4
```

Interpretation:

- `fill(person)` fills missing values in the specified column (`person` in this case) with the most recent non-missing value in that column.

Transform/manipulate data with “dplyr” To learn five key “dplyr” package functions that allow you to solve the vast majority of your data manipulation challenges:

- **filter()** function is used to subset a data frame, retaining all rows that satisfy your conditions.
- **arrange()**: orders the rows of a data frame by the values of selected columns.
- **select()**: (and optionally rename) variables in a data frame, using a concise mini-language that makes it easy to refer to variables based on their name e.g. a:f selects all columns from a on the left to f on the right) or type (e.g. where(is.numeric) selects all numeric columns).
- **mutate()**: creates new columns that are functions of existing variables. It can also modify (if the name is the same as an existing column) and delete columns (by setting their value to NULL). Collapse many values down to a single summary
- **summarise()**: It returns one row for each combination of grouping variables; if there are no grouping variables, the output will have a single row summarizing all observations in the input. It will contain one column for each grouping variable and one column for each of the summary statistics that you have specified
- **group_by()**: Most data operations are done on groups defined by variables. group_by() takes an existing tbl and converts it into a grouped tbl where operations are performed “by group”. ungroup() removes grouping.

Code Sample

```
library(dplyr)
library(nycflights13)
flights

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2       830           819
## 2  2013     1     1     533             529           4       850           830
## 3  2013     1     1     542             540           2       923           850
## 4  2013     1     1     544             545          -1      1004          1022
## 5  2013     1     1     554             600          -6       812           837
## 6  2013     1     1     554             558          -4       740           728
## 7  2013     1     1     555             600          -5       913           854
## 8  2013     1     1     557             600          -3       709           723
## 9  2013     1     1     557             600          -3       838           846
## 10 2013     1     1     558             600          -2       753           745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- load the flights dataset.

Filter

Code Sample

```
filter(flights, month == 1, day == 1)

## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 832 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- filters flights datasets for all months and day equals to 1

Code Sample

```
# filter jan1
jan1 <- filter(flights, month == 1, day == 1)
(jan1 <- filter(flights, month == 1, day == 1))

## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 832 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>

# filter dec25
dec25 <- filter(flights, day == 12, day == 25)
(dec25 <- filter(flights, day == 12, day == 25))
```



```
## # A tibble: 0 x 19
## # i 19 variables: year <int>, month <int>, day <int>, dep_time <int>,
## #   sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Code Sample *Interpretation:*

- filters January 1 for all flights data.
- filters December 25 for all flights data.

Code Sample

```
filter(flights, month = 1) #throws error
```

```
## Error in 'filter()':
## ! We detected a named input.
## i This usually means that you've used '=' instead of '=='.
## i Did you mean 'month == 1'?
```

```
filter(flights, month == 1) #works fine
```

```
## # A tibble: 27,004 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 26,994 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- `filter(flights, month = 1)` will throw an error because `month=1` doesn't filter but assigns the value 1 to month.
- `filter(flights, month == 1)` will filter the month 1.

Code Sample

```
# valid filter to fetch only rows where the month is either November December.
filter(flights, month == 11 | month == 12)

## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013    11     1       5           2359         6     352           345
## 2  2013    11     1      35           2250       105     123           2356
## 3  2013    11     1     455           500        -5     641           651
## 4  2013    11     1     539           545        -6     856           827
## 5  2013    11     1     542           545        -3     831           855
## 6  2013    11     1     549           600       -11     912           923
## 7  2013    11     1     550           600       -10     705           659
## 8  2013    11     1     554           600        -6     659           701
## 9  2013    11     1     554           600        -6     826           827
##10  2013    11     1     554           600        -6     749           751
## # i 55,393 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>

# returns all row since 11 | 12 computes to TRUE
filter(flights, month == 11 | 12)

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## 7  2013     1     1     555           600        -5     913           854
## 8  2013     1     1     557           600        -3     709           723
## 9  2013     1     1     557           600        -3     838           846
##10  2013     1     1     558           600        -2     753           745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>

# valid filter to fetch only rows where the month is either November December.
nov_dec <- filter(flights, month %in% c(11, 12))
```

Interpretation:

- `filter(flights, month == 11 | month == 12)` fetch only rows where the month is either November December.
- `filter(flights, month == 11 | 12)` returns all row since `11 | 12` computes to `TRUE`, even though syntactically it returns values but it won't gives us desired output.
- `nov_dec <- filter(flights, month %in% c(11, 12))` fetch only rows where the month is either November December.

Code Sample

```
# De Morgan's Law:
filter(flights, !(arr_delay > 120 | dep_delay > 120))

## # A tibble: 316,050 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## 7  2013     1     1     555           600        -5     913           854
## 8  2013     1     1     557           600        -3     709           723
## 9  2013     1     1     557           600        -3     838           846
## 10 2013     1     1     558           600        -2     753           745
## # i 316,040 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>

filter(flights, arr_delay <= 120, dep_delay <= 120)

## # A tibble: 316,050 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## 7  2013     1     1     555           600        -5     913           854
## 8  2013     1     1     557           600        -3     709           723
## 9  2013     1     1     557           600        -3     838           846
## 10 2013     1     1     558           600        -2     753           745
## # i 316,040 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- Both returns the same value: `filter(flights, !(arr_delay > 120 | dep_delay > 120))` filters all arrival and departure delayed which are not greater 120.
- `filter(flights, arr_delay <= 120, dep_delay <= 120)` filters all arrival and departure delayed which are less than or equal to 120.

Arrange Example

Code Sample

by default sorts in ascending order

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2       830           819
## 2  2013     1     1     533           529         4       850           830
## 3  2013     1     1     542           540         2       923           850
## 4  2013     1     1     544           545        -1      1004          1022
## 5  2013     1     1     554           600        -6       812           837
## 6  2013     1     1     554           558        -4       740           728
## 7  2013     1     1     555           600        -5       913           854
## 8  2013     1     1     557           600        -3       709           723
## 9  2013     1     1     557           600        -3       838           846
## 10 2013     1     1     558           600        -2       753           745
```

```
## # i 336,766 more rows
```

```
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

descending order

```
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     9     641           900      1301      1242          1530
## 2  2013     6    15    1432          1935      1137      1607          2120
## 3  2013     1    10    1121          1635      1126      1239          1810
## 4  2013     9    20    1139          1845      1014      1457          2210
## 5  2013     7    22     845          1600      1005      1044          1815
## 6  2013     4    10    1100          1900       960      1342          2211
## 7  2013     3    17    2321           810       911       135          1020
## 8  2013     6    27     959          1900       899      1236          2226
## 9  2013     7    22    2257           759       898       121          1026
## 10 2013    12     5     756          1700       896      1058          2020
```

```
## # i 336,766 more rows
```

```
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- `arrange(flights, year, month, day)` sorts year, month and day in ascending order.
- `arrange(flights, year, month, day)` sorts year, month and day in descending order.
- Missing values are always sorted at the end.

Select Example

Code Sample

```
# Select columns by name
```

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
```

```
##   year month   day
```

```
##   <int> <int> <int>
```

```
## 1  2013     1     1
```

```
## 2  2013     1     1
```

```
## 3  2013     1     1
```

```
## 4  2013     1     1
```

```
## 5  2013     1     1
```

```
## 6  2013     1     1
```

```
## 7  2013     1     1
```

```
## 8  2013     1     1
```

```
## 9  2013     1     1
```

```
## 10 2013     1     1
```

```
## # i 336,766 more rows
```

```
# Select all columns between year and day (inclusive)
```

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
```

```
##   year month   day
```

```
##   <int> <int> <int>
```

```
## 1  2013     1     1
```

```
## 2  2013     1     1
```

```
## 3  2013     1     1
```

```
## 4  2013     1     1
```

```
## 5  2013     1     1
```

```
## 6  2013     1     1
```

```
## 7  2013     1     1
```

```
## 8  2013     1     1
```

```
## 9  2013     1     1
```

```
## 10 2013     1     1
```

```
## # i 336,766 more rows
```

```
# Select all columns except those from year to day (inclusive)
```

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
```

```
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
```

```
##   <int>         <int>    <dbl>    <int>         <int>    <dbl> <chr>
```

```
## 1     517           515        2      830           819      11 UA
```

```
## 2     533           529        4      850           830      20 UA
```

```
## 3     542           540        2      923           850      33 AA
```

```
## 4     544           545       -1     1004          1022     -18 B6
```

```
## 5     554           600       -6      812           837     -25 DL
```

```
## 6     554           558       -4      740           728      12 UA
```

```
## 7     555           600       -5      913           854      19 B6
```

```
## 8     557           600       -3      709           723     -14 EV
```

```
## 9     557           600       -3      838           846      -8 B6
```

```
## 10      558      600      -2      753      745      8 AA
## # i 336,766 more rows
## # i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- `select(flights, year, month, day)` selects year, month and day from flights.
- `select(flights, year:day)` selects all columns between year and day (inclusive) from flights.
- `select(flights, year, month, day)` selects all columns except those from year to day (inclusive).

Other useful select functions

- `flights %>% select(starts_with('d'))` #matches names beginning with d
- `flights %>% select(ends_with('e'))` #matches names ending with e
- `flights %>% select(contains('d'))` #matches names containing with d
- `flights %>% select(matches("(.)\\1"))` #matches any variables that contain repeated characters
- `flights %>% select(num_range("dest",1:15))` #matches if dest contains any -range between 1 to 15

Code Sample

```
flights %>% select(starts_with('d')) #matches names beginning with d
```

```
## # A tibble: 336,776 x 5
##   day dep_time dep_delay dest distance
##   <int>   <int>   <dbl> <chr>   <dbl>
## 1     1     517         2 IAH     1400
## 2     1     533         4 IAH     1416
## 3     1     542         2 MIA     1089
## 4     1     544        -1 BQN     1576
## 5     1     554        -6 ATL      762
## 6     1     554        -4 ORD      719
## 7     1     555        -5 FLL     1065
## 8     1     557        -3 IAD      229
## 9     1     557        -3 MCO      944
## 10    1     558        -2 ORD      733
## # i 336,766 more rows
```

```
flights %>% select(ends_with('e')) #matches names ending with e
```

```
## # A tibble: 336,776 x 7
##   dep_time sched_dep_time arr_time sched_arr_time air_time distance minute
##   <int>       <int>       <int>       <int>       <dbl>   <dbl>   <dbl>
## 1     517         515         830         819        227    1400    15
## 2     533         529         850         830        227    1416    29
## 3     542         540         923         850        160    1089    40
## 4     544         545        1004        1022        183    1576    45
## 5     554         600         812         837        116     762     0
## 6     554         558         740         728        150     719    58
## 7     555         600         913         854        158    1065     0
```

```
## 8      557      600      709      723      53      229      0
## 9      557      600      838      846      140      944      0
## 10     558      600      753      745      138      733      0
## # i 336,766 more rows
```

```
flights %>% select(contains('d')) #matches names containing with d
```

```
## # A tibble: 336,776 x 8
##   day dep_time sched_dep_time dep_delay sched_arr_time arr_delay dest
##   <int>   <int>         <int>    <dbl>         <int>    <dbl> <chr>
## 1     1     517           515        2           819        11 IAH
## 2     1     533           529        4           830        20 IAH
## 3     1     542           540        2           850        33 MIA
## 4     1     544           545       -1          1022       -18 BQN
## 5     1     554           600       -6           837       -25 ATL
## 6     1     554           558       -4           728        12 ORD
## 7     1     555           600       -5           854        19 FLL
## 8     1     557           600       -3           723       -14 IAD
## 9     1     557           600       -3           846        -8 MCO
## 10    1     558           600       -2           745         8 ORD
## # i 336,766 more rows
## # i 1 more variable: distance <dbl>
```

```
flights %>% select(matches("(.)\\1")) #matches any variables that contain repeated characters
```

```
## # A tibble: 336,776 x 4
##   arr_time sched_arr_time arr_delay carrier
##   <int>         <int>    <dbl> <chr>
## 1     830           819        11 UA
## 2     850           830        20 UA
## 3     923           850        33 AA
## 4    1004          1022       -18 B6
## 5     812           837       -25 DL
## 6     740           728        12 UA
## 7     913           854        19 B6
## 8     709           723       -14 EV
## 9     838           846        -8 B6
## 10     753           745         8 AA
## # i 336,766 more rows
```

```
flights %>% select(num_range("dest",1:15)) #matches if dest contains any range between 1 to 15
```

```
## # A tibble: 336,776 x 0
```

Code Sample

```
# rename
rename(flights, tail_num = tailnum)

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```
##      <int> <int> <int>      <int>      <int>      <dbl>      <int>      <int>
## 1  2013      1      1      517      515          2      830      819
## 2  2013      1      1      533      529          4      850      830
## 3  2013      1      1      542      540          2      923      850
## 4  2013      1      1      544      545         -1     1004     1022
## 5  2013      1      1      554      600         -6      812      837
## 6  2013      1      1      554      558         -4      740      728
## 7  2013      1      1      555      600         -5      913      854
## 8  2013      1      1      557      600         -3      709      723
## 9  2013      1      1      557      600         -3      838      846
## 10 2013      1      1      558      600         -2      753      745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tail_num <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# move to the start of the data frame
select(flights, time_hour, air_time, everything())
```

```
## # A tibble: 336,776 x 19
##   time_hour      air_time year month   day dep_time sched_dep_time
##   <dtm>      <dbl> <int> <int> <int>   <int>      <int>
## 1 2013-01-01 05:00:00    227  2013     1     1     517          515
## 2 2013-01-01 05:00:00    227  2013     1     1     533          529
## 3 2013-01-01 05:00:00    160  2013     1     1     542          540
## 4 2013-01-01 05:00:00    183  2013     1     1     544          545
## 5 2013-01-01 06:00:00    116  2013     1     1     554          600
## 6 2013-01-01 05:00:00    150  2013     1     1     554          558
## 7 2013-01-01 06:00:00    158  2013     1     1     555          600
## 8 2013-01-01 06:00:00     53  2013     1     1     557          600
## 9 2013-01-01 06:00:00    140  2013     1     1     557          600
## 10 2013-01-01 06:00:00    138  2013     1     1     558          600
## # i 336,766 more rows
## # i 12 more variables: dep_delay <dbl>, arr_time <int>, sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, distance <dbl>, hour <dbl>, minute <dbl>
```

Interpretation:

- `select()` can be used to rename variables, but it's rarely useful because it drops all of the variables not explicitly.
- so using `rename()`, which is a variant of `select()` that keeps all the variables that aren't explicitly mentioned.
- Another option is to use `select()` in conjunction with the `everything()` helper.

Mutate: Example

Code Sample

```
flights_sml <- select(flights,
                      year:day,
                      ends_with("delay"),
```



```
distance,
air_time)
```

Interpretation:

- adding some variables in flights_sml from flights.

Code Sample

```
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  speed = distance / air_time * 60)

## # A tibble: 336,776 x 9
##   year month   day dep_delay arr_delay distance air_time gain speed
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>   <dbl> <dbl> <dbl>
## 1  2013     1     1         2         11     1400    227    -9   370.
## 2  2013     1     1         4         20     1416    227   -16   374.
## 3  2013     1     1         2         33     1089    160   -31   408.
## 4  2013     1     1        -1        -18     1576    183    17   517.
## 5  2013     1     1        -6        -25      762    116    19   394.
## 6  2013     1     1        -4         12      719    150   -16   288.
## 7  2013     1     1        -5         19     1065    158   -24   404.
## 8  2013     1     1        -3        -14      229     53    11   259.
## 9  2013     1     1        -3         -8      944    140     5   405.
## 10 2013     1     1        -2         8      733    138   -10   319.
## # i 336,766 more rows
```

Interpretation:

- two new columns, gain and speed, is added based on existing variables in the flights_sml dataset.

Code Sample

```
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours
)

## # A tibble: 336,776 x 10
##   year month   day dep_delay arr_delay distance air_time gain hours
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>   <dbl> <dbl>
## 1  2013     1     1         2         11     1400    227    -9  3.78
## 2  2013     1     1         4         20     1416    227   -16  3.78
## 3  2013     1     1         2         33     1089    160   -31  2.67
## 4  2013     1     1        -1        -18     1576    183    17  3.05
## 5  2013     1     1        -6        -25      762    116    19  1.93
## 6  2013     1     1        -4         12      719    150   -16  2.5
## 7  2013     1     1        -5         19     1065    158   -24  2.63
## 8  2013     1     1        -3        -14      229     53    11  0.883
```

```
## 9 2013 1 1 -3 -8 944 140 5 2.33
## 10 2013 1 1 -2 8 733 138 -10 2.3
## # i 336,766 more rows
## # i 1 more variable: gain_per_hour <dbl>
```

Interpretation:

- three new columns (gain, hours, and gain_per_hour) were added based on existing variables which are computed in the flights_sml dataset.

Code Sample

```
# transmute
# If you only want to keep the new variables, use transmute()
transmute(
  flights,
  gain = dep_delay - arr_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours
)
```

```
## # A tibble: 336,776 x 3
##   gain hours gain_per_hour
##   <dbl> <dbl>         <dbl>
## 1    -9 3.78          -2.38
## 2   -16 3.78          -4.23
## 3   -31 2.67         -11.6
## 4    17 3.05           5.57
## 5    19 1.93           9.83
## 6   -16 2.5           -6.4
## 7   -24 2.63          -9.11
## 8    11 0.883         12.5
## 9     5 2.33           2.14
## 10  -10 2.3          -4.35
## # i 336,766 more rows
```

Interpretation:

- three new variables (gain, hours, and gain_per_hour) based on existing variables in the flights dataset while discarding all other variables.

Code Sample

```
transmute(flights,
  dep_time,
  hour = dep_time %/% 100,
  minute = dep_time %% 100)

## # A tibble: 336,776 x 3
##   dep_time hour minute
##   <int> <dbl> <dbl>
```

```
## 1      517      5      17
## 2      533      5      33
## 3      542      5      42
## 4      544      5      44
## 5      554      5      54
## 6      554      5      54
## 7      555      5      55
## 8      557      5      57
## 9      557      5      57
## 10     558      5      58
## # i 336,766 more rows
```

Interpretation:

- this code creates two new variables (hour and minute) from the dep_time variable in the flights dataset, while still retaining the original dep_time variable in the output and discard all other columns.

Code Sample

```
# Select columns by name
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # i 336,766 more rows
```

```
# Select all columns between year and day (inclusive)
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # i 336,766 more rows
```

```
# Select all columns except those from year to day (inclusive)
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>      <int>      <dbl>   <int>      <int>      <dbl> <chr>
## 1      517          515         2     830          819        11 UA
## 2      533          529         4     850          830        20 UA
## 3      542          540         2     923          850        33 AA
## 4      544          545        -1    1004         1022       -18 B6
## 5      554          600        -6     812          837       -25 DL
## 6      554          558        -4     740          728        12 UA
## 7      555          600        -5     913          854        19 B6
## 8      557          600        -3     709          723       -14 EV
## 9      557          600        -3     838          846        -8 B6
## 10     558          600        -2     753          745         8 AA
## # i 336,766 more rows
## # i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- `select(flights, year, month, day)` selects year, month and day from flights.
- `select(flights, year:day)` selects all columns between year and day (inclusive) from flights.
- `select(flights, year, month, day)` selects all columns except those from year to day (inclusive).

Summarise Example

Code Sample

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))

## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

Interpretation:

- computes the mean departure delay for the flights dataset, removing any missing values in the `dep_delay` column.

Code Sample

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))

## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

Interpretation:

- computes the mean departure delay for the flights dataset, removing any missing values in the dep_delay column.

Code Sample

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))

## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

Interpretation:

- computes the mean departure delay for the flights dataset, removing any missing values in the dep_delay column.

Code Sample

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))

## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

Interpretation:

- computes the mean departure delay for the flights dataset, removing any missing values in the dep_delay column.

Code Sample

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))

## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

Interpretation:

- computes the mean departure delay for the flights dataset, removing any missing values in the dep_delay column.

Code Sample

```
by_day <- group_by(flights, year, month, day)
summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))

## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.

## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # i 355 more rows
```

Interpretation:

- computes the mean departure delay for each day in the dataset by grouping the data by year, month, and day, and then summarizing the grouped data by average delayed time.

Code Sample

```
delays <- flights %>%
  group_by(dest) %>%
  summarise(
    count = n(),
    dist = mean(distance, na.rm = TRUE),
    delay = mean(arr_delay, na.rm = TRUE)
  ) %>%
  filter(count > 20, dest != "HNL")
delays

## # A tibble: 96 x 4
##   dest count dist delay
##   <chr> <int> <dbl> <dbl>
## 1 ABQ    254 1826   4.38
## 2 ACK    265  199   4.85
## 3 ALB    439  143  14.4
## 4 ATL  17215  757.  11.3
## 5 AUS   2439 1514.   6.02
## 6 AVL    275  584.   8.00
## 7 BDL    443  116   7.05
```

```
## 8 BGR      375  378   8.03
## 9 BHM      297  866. 16.9
## 10 BNA     6333  758. 11.8
## # i 86 more rows
```

Interpretation:

- computes summary statistics on flight delays, including the count of flights, the average distance, and the average arrival delay, grouped by destination. It then filters the results to exclude destinations with greater than 20 flights and the destination of Honolulu.

Code Sample

```
flights %>%
  group_by(year, month, day) %>%
  summarise(mean = mean(dep_delay))

## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.

## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day mean
##   <int> <int> <int> <dbl>
## 1  2013     1     1    NA
## 2  2013     1     2    NA
## 3  2013     1     3    NA
## 4  2013     1     4    NA
## 5  2013     1     5    NA
## 6  2013     1     6    NA
## 7  2013     1     7    NA
## 8  2013     1     8    NA
## 9  2013     1     9    NA
## 10 2013     1    10    NA
## # i 355 more rows
```

Interpretation:

- calculates the mean departure delay for each day by grouping the data by year, month, and day, and then summarising the grouped data.

Code Sample

```
flights %>%
  group_by(year, month, day) %>%
  summarise(mean = mean(dep_delay, na.rm = TRUE))

## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day mean
##   <int> <int> <int> <dbl>
## 1  2013     1     1 11.5
## 2  2013     1     2 13.9
## 3  2013     1     3 11.0
## 4  2013     1     4  8.95
## 5  2013     1     5  5.73
## 6  2013     1     6  7.15
## 7  2013     1     7  5.42
## 8  2013     1     8  2.55
## 9  2013     1     9  2.28
## 10 2013     1    10  2.84
## # i 355 more rows
```

Interpretation:

- calculates the mean departure delay for each day by grouping the data by year, month, and day, and then summarizing the grouped data after removing null from the data in the dep_delay column while computing the mean.

Code Sample

```
not_cancelled <- flights %>%
  filter(!is.na(dep_delay), !is.na(arr_delay))
```

Interpretation:

- not_cancelled contains only flights that were not cancelled, based on the absence of missing values in the departure delay and arrival delay columns.
- ! negates the TRUE value to FALSE and vice versa.

Code Sample

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(mean =
    mean(dep_delay))
```

```
## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day mean
##   <int> <int> <int> <dbl>
## 1  2013     1     1 11.4
## 2  2013     1     2 13.7
## 3  2013     1     3 10.9
## 4  2013     1     4  8.97
```



```
## 5 2013 1 5 5.73
## 6 2013 1 6 7.15
## 7 2013 1 7 5.42
## 8 2013 1 8 2.56
## 9 2013 1 9 2.30
## 10 2013 1 10 2.84
## # i 355 more rows
```

Interpretation:

- calculates the mean departure delay for each day in the not_cancelled dataset by grouping the data by year, month, and day, and then summarizing the grouped data.

Code Sample

```
delays <- not_cancelled %>%
  group_by(tailnum) %>%
  summarise(
    delay = mean(arr_delay)
  )
delays
```


```
## # A tibble: 4,037 x 2
##   tailnum delay
##   <chr>   <dbl>
## 1 D942DN  31.5
## 2 NOEGMQ   9.98
## 3 N10156  12.7
## 4 N102UW   2.94
## 5 N103US  -6.93
## 6 N104UW   1.80
## 7 N10575  20.7
## 8 N105UW  -0.267
## 9 N107US  -5.73
## 10 N108UW -1.25
## # i 4,027 more rows
```

Interpretation:

- calculates the average arrival delay for each planes in the not_cancelled dataset by grouping the data by tailnum and then summarizing the grouped data.

Code Sample

```
hist(delays$delay)
```

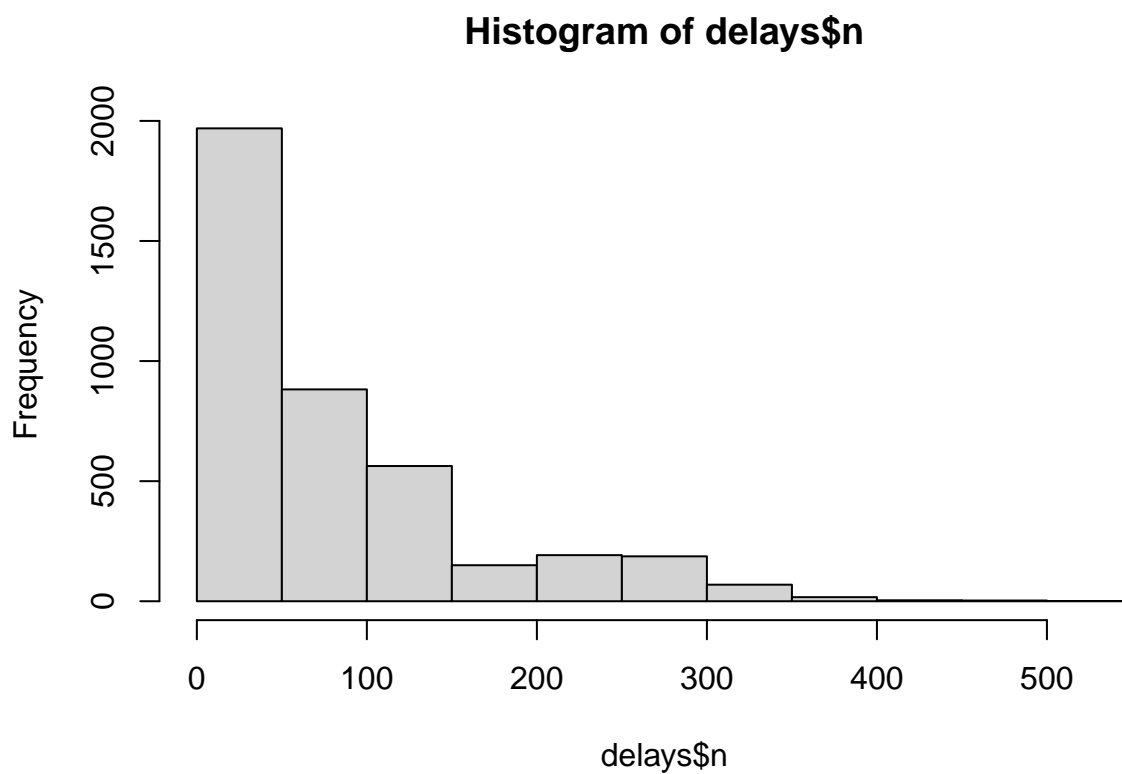


Interpretation:

- creates a histogram of the average arrival delays (delay) column.

Code Sample

```
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarise(  
    delay = mean(arr_delay,  
                 na.rm = TRUE),  
    n = n()  
  )  
hist(delays$n)
```




Interpretation:

- creates a histogram of the counts (n) of flights for each tailnum.

Code Sample

```
plot(delays$n, delays$delay)
```



Interpretation:

- creates a scatter plot showing the relationship between the count of flights and the average arrival delay for each aircraft.
- even though creating histogram and scatter plot is good idea but if we can't make sense or interpret it, it will be useless in analysis.

Code Sample

```
# When do the first and last flights leave each day
not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(
    first = min(dep_time),
    last = max(dep_time)
  )
```

```
## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 365 x 5
## # Groups:   year, month [12]
##   year month   day first last
##   <int> <int> <int> <int> <int>
## 1  2013     1     1   517 2356
```

```
## 2 2013 1 2 42 2354
## 3 2013 1 3 32 2349
## 4 2013 1 4 25 2358
## 5 2013 1 5 14 2357
## 6 2013 1 6 16 2355
## 7 2013 1 7 49 2359
## 8 2013 1 8 454 2351
## 9 2013 1 9 2 2252
## 10 2013 1 10 3 2320
## # i 355 more rows
```

Interpretation:

- calculates the first and last departure times for each day in the not_cancelled dataset by grouping the data by year, month, and day, and then summarizing the grouped data.

Code Sample

```
# Why is distance to some destinations more variable than to others?
not_cancelled %>%
  group_by(dest) %>%
  summarise(distance_sd =
    sd(distance)) %>%
  arrange(desc(distance_sd))

## # A tibble: 104 x 2
##   dest distance_sd
##   <chr>         <dbl>
## 1 EGE          10.5
## 2 SAN          10.4
## 3 SFO          10.2
## 4 HNL          10.0
## 5 SEA          9.98
## 6 LAS          9.91
## 7 PDX          9.87
## 8 PHX          9.86
## 9 LAX          9.66
## 10 IND         9.46
## # i 94 more rows
```

Interpretation:

- provides information about the variability of distances for each destination in the not_cancelled dataset, with destinations having higher standard deviations of distances appearing first in the resulting dataset.

Code Sample

```
# Which destinations have the most carriers?
not_cancelled %>%
  group_by(dest) %>%
  summarise(carriers =
    n_distinct(carrier)) %>%
  arrange(desc(carriers))
```

```
## # A tibble: 104 x 2
##   dest carriers
##   <chr>     <int>
## 1 ATL         7
## 2 BOS         7
## 3 CLT         7
## 4 ORD         7
## 5 TPA         7
## 6 AUS         6
## 7 DCA         6
## 8 DTW         6
## 9 IAD         6
## 10 MSP        6
## # i 94 more rows
```

Interpretation:

- summarizes the destination by most number of carriers.

Code Sample

```
# How many flights left before 5am? (these usually indicate delayed flights from the previous day)
not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(n_early =
    sum(dep_time < 500))

## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day n_early
##   <int> <int> <int>   <int>
## 1  2013     1     1         0
## 2  2013     1     2         3
## 3  2013     1     3         4
## 4  2013     1     4         3
## 5  2013     1     5         3
## 6  2013     1     6         2
## 7  2013     1     7         2
## 8  2013     1     8         1
## 9  2013     1     9         3
## 10 2013     1    10         3
## # i 355 more rows
```

Interpretation:

- summarizes the number of flights departing before 5:00 AM for each day in the not_cancelled dataset, grouped by year, month, and day.

Code Sample

```
# What proportion of flights are delayed by more than an hour?
not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(hour_prop =
    mean(arr_delay > 60))

## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.

## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day hour_prop
##   <int> <int> <int>     <dbl>
## 1  2013     1     1    0.0722
## 2  2013     1     2    0.0851
## 3  2013     1     3    0.0567
## 4  2013     1     4    0.0396
## 5  2013     1     5    0.0349
## 6  2013     1     6    0.0470
## 7  2013     1     7    0.0333
## 8  2013     1     8    0.0213
## 9  2013     1     9    0.0202
## 10 2013     1    10    0.0183
## # i 355 more rows
```

Interpretation:

- summarizes the proportion of flights with arrival delays exceeding 60 minutes for each day in the not_cancelled dataset, grouped by year, month, and day.

Code Sample

```
#Find all groups bigger than a threshold:
popular_dests <- flights %>%
  group_by(dest) %>%
  filter(n() > 365)
popular_dests

## # A tibble: 332,577 x 19
## # Groups:   dest [77]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>     <int>         <int>
## 1  2013     1     1     517           515         2       830           819
## 2  2013     1     1     533           529         4       850           830
## 3  2013     1     1     542           540         2       923           850
## 4  2013     1     1     544           545        -1      1004          1022
## 5  2013     1     1     554           600        -6       812           837
## 6  2013     1     1     554           558        -4       740           728
## 7  2013     1     1     555           600        -5       913           854
## 8  2013     1     1     557           600        -3       709           723
```

```
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # i 332,567 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
head(popular_dests$dest)
```

```
## [1] "IAH" "IAH" "MIA" "BQN" "ATL" "ORD"
```

```
tail(popular_dests$dest)
```

```
## [1] "BNA" "DCA" "SYR" "BNA" "CLE" "RDU"
```

Interpretation:

- summarizes the destinations that are popular based on having more than 365 flights in the flights dataset.
- head() and tail() functions are then used to display the first and last few destinations in the popular_dests dataset, respectively.

Slice function

Code Sample

```
#What will happen?
flights %>% slice(1L)

## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- slice the first row of the flights dataset.

Code Sample

```
#What will happen?
flights %>% slice(n())

## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     9    30      NA           840        NA     NA           1020
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation: - slice the last row of the flights dataset.

Code Sample

```
#What will happen?
flights %>% slice(5:n())

## # A tibble: 336,772 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     554             600          -6     812             837
## 2  2013     1     1     554             558          -4     740             728
## 3  2013     1     1     555             600          -5     913             854
## 4  2013     1     1     557             600          -3     709             723
## 5  2013     1     1     557             600          -3     838             846
## 6  2013     1     1     558             600          -2     753             745
## 7  2013     1     1     558             600          -2     849             851
## 8  2013     1     1     558             600          -2     853             856
## 9  2013     1     1     558             600          -2     924             917
## 10 2013     1     1     558             600          -2     923             937
## # i 336,762 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- will slice rows from the 5th row to the last row of the flights dataset.

Code Sample

```
#What will happen?
slice(flights,-(1:4))

## # A tibble: 336,772 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     554             600          -6     812             837
## 2  2013     1     1     554             558          -4     740             728
## 3  2013     1     1     555             600          -5     913             854
## 4  2013     1     1     557             600          -3     709             723
## 5  2013     1     1     557             600          -3     838             846
## 6  2013     1     1     558             600          -2     753             745
## 7  2013     1     1     558             600          -2     849             851
## 8  2013     1     1     558             600          -2     853             856
## 9  2013     1     1     558             600          -2     924             917
## 10 2013     1     1     558             600          -2     923             937
## # i 336,762 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- will remove the first four rows from the flights dataset.

Code Sample

```
#What will happen?
flights %>% slice_sample(n=5)

## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     4    11    1610           1529        41     1729           1648
## 2  2013    10    25     626           600        26      733           715
## 3  2013     2    28    1125           1115        10     1229           1255
## 4  2013     9    28     822           822         0     1248           1215
## 5  2013     5    27    1221           1230        -9     1314           1338
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- will randomly sample 5 rows from the flights dataset without replacement.

Code Sample

```
#What will happen?
flights %>% slice_sample(n=5, replace = TRUE)

## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     4    24    1859           1815        44     2115           1948
## 2  2013    11    15    1947           1955        -8     2147           2207
## 3  2013     5    18     843           820        23      951           930
## 4  2013     3    22     815           820        -5      928           944
## 5  2013    12     1    1603           1600         3     1807           1824
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- will randomly sample 5 rows from the flights dataset with replacement, meaning the same row may be selected more than once.

Code Sample

```
#What will happen?
set.seed(123)
train_data <- flights %>% slice_sample(prop=0.8)
train_data
```

```
## # A tibble: 269,420 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     4    26     730             655        35     901           820
## 2  2013     2    26    1714            1720       -6    2031          2040
## 3  2013     2    15    1442            1445       -3    1634          1647
## 4  2013     3    27    1444            1445       -1    1556          1604
## 5  2013     6     5    1428            1430       -2    1537          1555
## 6  2013     2    16     613             600        13     731           735
## 7  2013     5     1     848             850       -2     947          1014
## 8  2013    10    20    1918            1855        23    2132          2130
## 9  2013    11    11     825             835       -10     942          1000
## 10 2013     8    31     656             700        -4     904           929
## # i 269,410 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- sets the random seed to ensure reproducibility of random sampling.
- `train_data <- flights %>% slice_sample(prop = 0.8)` creates a training dataset (`train_data`) by randomly sampling 80% of the rows from the `flights` dataset without replacement.

Code Sample

```
test_data <- flights %>% slice_sample(prop=0.2)
test_data

## # A tibble: 67,355 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     8    23    1605            1559         6    1700          1721
## 2  2013     5     2    1222            1205        17    1444          1425
## 3  2013     4     7    1957            2000       -3    2204          2208
## 4  2013     4    16     639             640       -1     837           851
## 5  2013     4     2    1132            1135       -3    1311          1255
## 6  2013     8     7    1006             959         7    1116          1114
## 7  2013     9    26    1826            1829       -3    2014          2033
## 8  2013     5    19    1945            1925        20    2252          2250
## 9  2013     7    15    1654            1659       -5    1843          1905
## 10 2013     4    13     741             745       -4     959          1012
## # i 67,345 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interpretation:

- sets the random seed to ensure reproducibility of random sampling.
- `test_data <- flights %>% slice_sample(prop=0.2)` creates a testing dataset (`test_data`) by randomly sampling 20% of the rows from the `flights` dataset without replacement.