

FDS-A2: Data Wrangling

Suman Paudel

Task 1

```
library(tidyverse)
library(magrittr)
library(reticulate)
library(scales)
use_python("C:\\Python311\\python.exe")
```

Load the Data

```
# load the dataset
auto_data <- read.csv("automobile.data", header = FALSE, na.strings = "?")
```

Basic Data Exploration

```
# Basic Data exploration
# Check dimensions
dim(auto_data)
```

```
## [1] 205 26
```

```
# View first few rows
head(auto_data)
```

```
##   V1  V2      V3  V4  V5  V6      V7  V8    V9  V10  V11  V12  V13
## 1  3  NA alfa-romero gas std two convertible rwd front 88.6 168.8 64.1 48.8
## 2  3  NA alfa-romero gas std two convertible rwd front 88.6 168.8 64.1 48.8
## 3  1  NA alfa-romero gas std two hatchback rwd front 94.5 171.2 65.5 52.4
## 4  2 164      audi gas std four      sedan fwd front 99.8 176.6 66.2 54.3
## 5  2 164      audi gas std four      sedan 4wd front 99.4 176.6 66.4 54.3
## 6  2  NA      audi gas std two      sedan fwd front 99.8 177.3 66.3 53.1
##   V14  V15  V16 V17  V18  V19  V20  V21  V22  V23  V24  V25  V26
## 1 2548 dohc four 130 mpfi 3.47 2.68 9.0 111 5000 21 27 13495
## 2 2548 dohc four 130 mpfi 3.47 2.68 9.0 111 5000 21 27 16500
## 3 2823 ohcv six 152 mpfi 2.68 3.47 9.0 154 5000 19 26 16500
## 4 2337 ohc four 109 mpfi 3.19 3.40 10.0 102 5500 24 30 13950
## 5 2824 ohc five 136 mpfi 3.19 3.40 8.0 115 5500 18 22 17450
## 6 2507 ohc five 136 mpfi 3.19 3.40 8.5 110 5500 19 25 15250
```

```
# Check data types
str(auto_data)
```

```
## 'data.frame':    205 obs. of  26 variables:
## $ V1 : int  3 3 1 2 2 2 1 1 1 0 ...
## $ V2 : int  NA NA NA 164 164 NA 158 NA 158 NA ...
## $ V3 : chr  "alfa-romero" "alfa-romero" "alfa-romero" "audi" ...
## $ V4 : chr  "gas" "gas" "gas" "gas" ...
## $ V5 : chr  "std" "std" "std" "std" ...
## $ V6 : chr  "two" "two" "two" "four" ...
## $ V7 : chr  "convertible" "convertible" "hatchback" "sedan" ...
## $ V8 : chr  "rwd" "rwd" "rwd" "fwd" ...
## $ V9 : chr  "front" "front" "front" "front" ...
## $ V10: num  88.6 88.6 94.5 99.8 99.4 ...
## $ V11: num  169 169 171 177 177 ...
## $ V12: num  64.1 64.1 65.5 66.2 66.4 66.3 71.4 71.4 71.4 67.9 ...
## $ V13: num  48.8 48.8 52.4 54.3 54.3 53.1 55.7 55.7 55.9 52 ...
## $ V14: int  2548 2548 2823 2337 2824 2507 2844 2954 3086 3053 ...
## $ V15: chr  "dohc" "dohc" "ohcv" "ohc" ...
## $ V16: chr  "four" "four" "six" "four" ...
## $ V17: int  130 130 152 109 136 136 136 136 131 131 ...
## $ V18: chr  "mpfi" "mpfi" "mpfi" "mpfi" ...
## $ V19: num  3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...
## $ V20: num  2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...
## $ V21: num  9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
## $ V22: int  111 111 154 102 115 110 110 110 140 160 ...
## $ V23: int  5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...
## $ V24: int  21 21 19 24 18 19 19 19 17 16 ...
## $ V25: int  27 27 26 30 22 25 25 25 20 22 ...
## $ V26: int  13495 16500 16500 13950 17450 15250 17710 18920 23875 NA ...
```

```
# Summary statistics
summary(auto_data)
```

```
##           V1           V2           V3           V4
## Min.      :-2.0000  Min.   : 65  Length:205  Length:205
## 1st Qu.:  0.0000  1st Qu.: 94  Class :character  Class :character
## Median :  1.0000  Median :115  Mode  :character  Mode  :character
## Mean      : 0.8341  Mean   :122
## 3rd Qu.:  2.0000  3rd Qu.:150
## Max.      : 3.0000  Max.   :256
##           NA's      :41
##           V5           V6           V7           V8
## Length:205  Length:205  Length:205  Length:205
## Class :character  Class :character  Class :character  Class :character
## Mode  :character  Mode  :character  Mode  :character  Mode  :character
##
##
##
##           V9           V10          V11          V12
## Length:205  Min.      : 86.60  Min.      :141.1  Min.      :60.30
## Class :character  1st Qu.: 94.50  1st Qu.:166.3  1st Qu.:64.10
## Mode  :character  Median : 97.00  Median :173.2  Median :65.50
##           Mean      : 98.76  Mean      :174.0  Mean      :65.91
##           3rd Qu.:102.40  3rd Qu.:183.1  3rd Qu.:66.90
##           Max.      :120.90  Max.      :208.1  Max.      :72.30
```

```
##
##      V13      V14      V15      V16
## Min.   :47.80  Min.   :1488  Length:205  Length:205
## 1st Qu.:52.00  1st Qu.:2145  Class :character  Class :character
## Median :54.10  Median :2414  Mode  :character  Mode  :character
## Mean   :53.72  Mean   :2556
## 3rd Qu.:55.50  3rd Qu.:2935
## Max.   :59.80  Max.   :4066
##
##      V17      V18      V19      V20
## Min.   : 61.0  Length:205  Min.   :2.54  Min.   :2.070
## 1st Qu.: 97.0  Class :character  1st Qu.:3.15  1st Qu.:3.110
## Median :120.0  Mode  :character  Median :3.31  Median :3.290
## Mean   :126.9  Mean   :3.33  Mean   :3.255
## 3rd Qu.:141.0  3rd Qu.:3.59  3rd Qu.:3.410
## Max.   :326.0  Max.   :3.94  Max.   :4.170
##              NA's :4  NA's :4
##      V21      V22      V23      V24      V25
## Min.   : 7.00  Min.   : 48.0  Min.   :4150  Min.   :13.00  Min.   :16.00
## 1st Qu.: 8.60  1st Qu.: 70.0  1st Qu.:4800  1st Qu.:19.00  1st Qu.:25.00
## Median : 9.00  Median : 95.0  Median :5200  Median :24.00  Median :30.00
## Mean   :10.14  Mean   :104.3  Mean   :5125  Mean   :25.22  Mean   :30.75
## 3rd Qu.: 9.40  3rd Qu.:116.0  3rd Qu.:5500  3rd Qu.:30.00  3rd Qu.:34.00
## Max.   :23.00  Max.   :288.0  Max.   :6600  Max.   :49.00  Max.   :54.00
##              NA's :2  NA's :2
##      V26
## Min.   : 5118
## 1st Qu.: 7775
## Median :10295
## Mean   :13207
## 3rd Qu.:16500
## Max.   :45400
## NA's   :4
```

Data Wrangling

```
# there was no header in the dataset I had to change the column names
# change the names of columns of the auto_data dataframe
colnames(auto_data) <- c("symboling", "normalized_losses", "make", "fuel_type",
  "aspiration", "num_doors", "body_style",
  "drive_wheels", "engine_location", "wheel_base",
  "length", "width", "height", "curb_weight",
  "engine_type", "num_cylinders", "engine_size",
  "fuel_system", "bore", "stroke", "compression_ratio",
  "horsepower", "peak_rpm", "city_mpg",
  "highway_mpg", "price")

colnames(auto_data)

## [1] "symboling"      "normalized_losses" "make"
## [4] "fuel_type"      "aspiration"        "num_doors"
## [7] "body_style"     "drive_wheels"      "engine_location"
## [10] "wheel_base"     "length"            "width"
## [13] "height"         "curb_weight"       "engine_type"
```

```
## [16] "num_cylinders"      "engine_size"      "fuel_system"
## [19] "bore"               "stroke"           "compression_ratio"
## [22] "horsepower"        "peak_rpm"         "city_mpg"
## [25] "highway_mpg"       "price"
```

```
# Check for missing values
colSums(is.na(auto_data))
```

```
##      symboling normalized_losses      make      fuel_type
##      0              41              0              0
##      aspiration      num_doors      body_style      drive_wheels
##      0              2              0              0
##      engine_location      wheel_base      length      width
##      0              0              0              0
##      height      curb_weight      engine_type      num_cylinders
##      0              0              0              0
##      engine_size      fuel_system      bore      stroke
##      0              0              4              4
##      compression_ratio      horsepower      peak_rpm      city_mpg
##      0              2              2              0
##      highway_mpg      price
##      0              4
```

```
# Impute the missing values
```

```
auto_data <- auto_data %>%
  mutate(
    num_doors = ifelse(is.na(num_doors), names(which.max(table(auto_data$num_doors))), num_doors),
    bore = ifelse(is.na(bore), median(bore, na.rm = TRUE), bore),
    stroke = ifelse(is.na(stroke), median(stroke, na.rm = TRUE), stroke),
    horsepower = ifelse(is.na(horsepower), median(horsepower, na.rm = TRUE), horsepower),
    peak_rpm = ifelse(is.na(peak_rpm), median(peak_rpm, na.rm = TRUE), peak_rpm),
    price = ifelse(is.na(price), median(price, na.rm = TRUE), price)
  )
```

```
head(auto_data)
```

```
##      symboling normalized_losses      make fuel_type aspiration num_doors
## 1          3              NA alfa-romero      gas      std      two
## 2          3              NA alfa-romero      gas      std      two
## 3          1              NA alfa-romero      gas      std      two
## 4          2             164      audi      gas      std      four
## 5          2             164      audi      gas      std      four
## 6          2              NA      audi      gas      std      two
##      body_style drive_wheels engine_location wheel_base length width height
## 1 convertible      rwd      front      88.6 168.8 64.1 48.8
## 2 convertible      rwd      front      88.6 168.8 64.1 48.8
## 3  hatchback      rwd      front      94.5 171.2 65.5 52.4
## 4      sedan      fwd      front      99.8 176.6 66.2 54.3
## 5      sedan      4wd      front      99.4 176.6 66.4 54.3
## 6      sedan      fwd      front      99.8 177.3 66.3 53.1
##      curb_weight engine_type num_cylinders engine_size fuel_system bore stroke
## 1          2548      dohc      four      130      mpfi 3.47 2.68
## 2          2548      dohc      four      130      mpfi 3.47 2.68
```

```
## 3      2823      ohcv      six      152      mpfi 2.68  3.47
## 4      2337      ohc      four     109      mpfi 3.19  3.40
## 5      2824      ohc      five     136      mpfi 3.19  3.40
## 6      2507      ohc      five     136      mpfi 3.19  3.40
## compression_ratio horsepower peak_rpm city_mpg highway_mpg price
## 1          9.0          111    5000      21      27 13495
## 2          9.0          111    5000      21      27 16500
## 3          9.0          154    5000      19      26 16500
## 4         10.0          102    5500      24      30 13950
## 5          8.0          115    5500      18      22 17450
## 6          8.5          110    5500      19      25 15250
```

```
colSums(is.na(auto_data))
```

```
##      symboling normalized_losses      make      fuel_type
##          0          41          0          0
##      aspiration      num_doors      body_style      drive_wheels
##          0          0          0          0
##      engine_location      wheel_base      length      width
##          0          0          0          0
##          height      curb_weight      engine_type      num_cylinders
##          0          0          0          0
##      engine_size      fuel_system      bore      stroke
##          0          0          0          0
##      compression_ratio      horsepower      peak_rpm      city_mpg
##          0          0          0          0
##      highway_mpg      price
##          0          0
```

```
# for categorical value
```

```
auto_data <- auto_data %>%
```

```
  mutate(normalized_losses = replace_na(normalized_losses, mean(normalized_losses, na.rm = TRUE)))
```

```
auto_data$num_doors[is.na(auto_data$num_doors)] <- names(which.max(table(auto_data$num_doors)))
```

```
# Again check if there are null values.
```


```
colSums(is.na(auto_data))
```

```
##      symboling normalized_losses      make      fuel_type
##          0          0          0          0
##      aspiration      num_doors      body_style      drive_wheels
##          0          0          0          0
##      engine_location      wheel_base      length      width
##          0          0          0          0
##          height      curb_weight      engine_type      num_cylinders
##          0          0          0          0
##      engine_size      fuel_system      bore      stroke
##          0          0          0          0
##      compression_ratio      horsepower      peak_rpm      city_mpg
##          0          0          0          0
##      highway_mpg      price
##          0          0
```

```
# saving data in csv so that I can use the visualization from python later on
auto_data_output <- auto_data %>% select_if(is.numeric)
# Save the transformed data to a CSV file
write.csv(auto_data_output, "transformed_data.csv", row.names = FALSE)
```


Exploratory Data Analysis

```
# relationship between length and width of a automobile
plot(auto_data$length, auto_data$width, main='Lenght vs Width', xlab = 'Length', ylab = 'Width')
```




```
# price by make
make_price <- as_tibble(aggregate(price ~ make, auto_data, mean))
make_price$price <- as.integer(make_price$price)

ggplot(make_price, aes(x = make, y = price)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  geom_text(aes(label = price), vjust = -0.5, size = 2.5) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 40, vjust = 0.5, hjust = 1),
    plot.title = element_text(hjust = 0.5))+
  ylim(0,40000)+
  labs(x = "Make", y = "Mean Price", title = "Make by Price")
```



```
# relationship between horsepower by number of cylinders
hp_cyl<- as_tibble(aggregate(horsepower ~ num_cylinders, auto_data, median))
ggplot(hp_cyl, aes(x = horsepower, y = num_cylinders)) +
  geom_bar(stat = "identity", fill = "#996498") +
  geom_text(aes(label = horsepower), vjust = 0.5, hjust = -0.2) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1),
    plot.title = element_text(hjust = 0.5))+
  labs(x = "Horse Power", y = "Number of Cylinders", title = "Horsepower by Number of Cylinders")
```




```
# cylinder vs mpg (city, highway)
mpg_summary <- auto_data %>%
  group_by(num_cylinders) %>%
  summarize(mean_city_mpg = round(mean(city_mpg)), mean_highway_mpg = round(mean(highway_mpg)))
mpg_summary

## # A tibble: 7 x 3
##   num_cylinders mean_city_mpg mean_highway_mpg
##   <chr>          <dbl>          <dbl>
## 1 eight          15            19
## 2 five           20            24
## 3 four           27            33
## 4 six            18            24
## 5 three          47            53
## 6 twelve         13            17
## 7 two            17            23

mpg_summary %>%
  pivot_longer(c(mean_city_mpg, mean_highway_mpg),
    names_to = "mpg_type",
    values_to = "mean_mpg") %>%
  ggplot(aes(x = num_cylinders, y = mean_mpg, fill = mpg_type)) +
  geom_col(position = "dodge") +
  geom_text(aes(label = round(mean_mpg, 1)), position = position_dodge(width = 1), vjust = -0.5) +
  labs(x = "Number of Cylinders", y = "Mean MPG", fill = "MPG Type") +
```

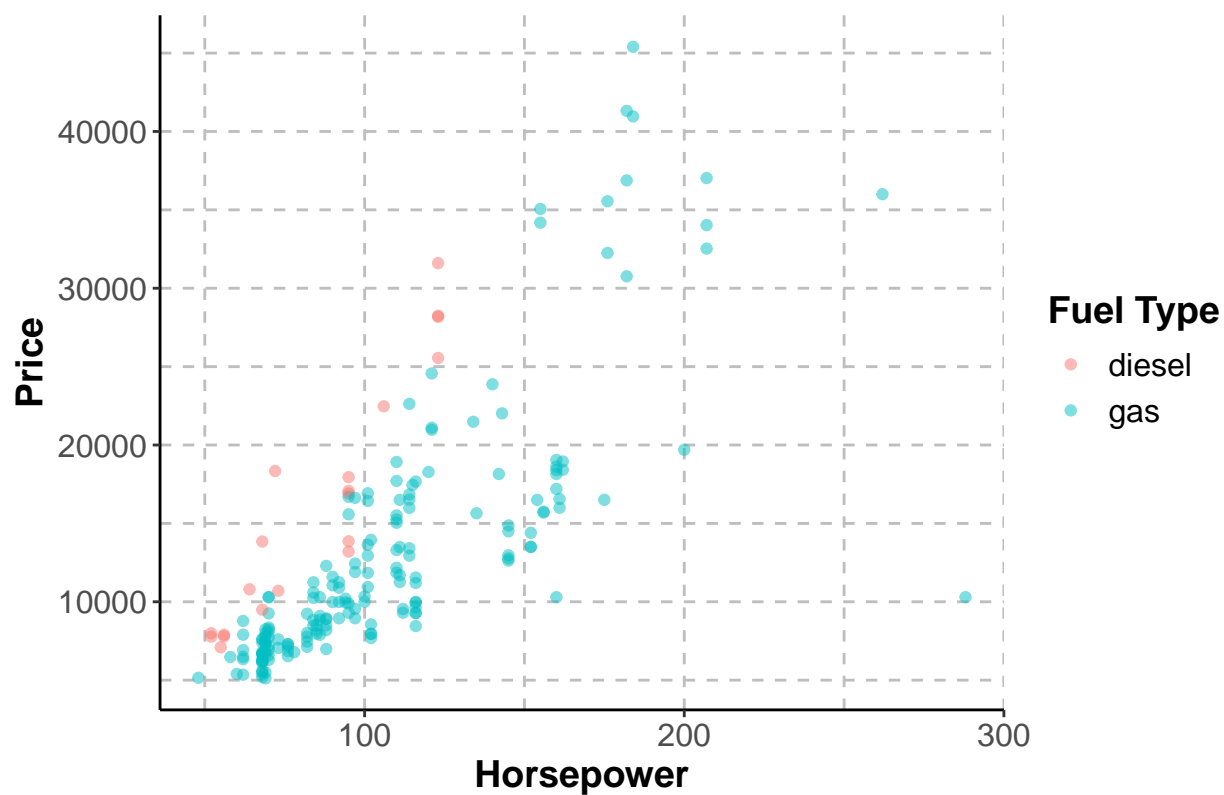


```
theme_classic() +
ggtitle("Relationship between Cylinders and Mileage")
```




```
# relationship between horsepower and price
ggplot(auto_data, aes(x = horsepower, y = price, color = fuel_type)) +
  geom_point(alpha = 0.5) +
  labs(x = "Horsepower", y = "Price", color = "Fuel Type") +
  theme_classic() +
  theme(panel.grid.major = element_line(size = 0.5, linetype = "dashed", color = "gray"),
        panel.grid.minor = element_line(size = 0.5, linetype = "dashed", color = "gray"),
        axis.text = element_text(size = 12),
        axis.title = element_text(size = 14, face = "bold"),
        plot.title = element_text(size = 16, face = "bold", hjust = 0.5),
        legend.text = element_text(size = 12),
        legend.title = element_text(size = 14, face = "bold")) +
  ggtitle("Relationship between Horsepower and Price")
```

Relationship between Horsepower and Price



```
# relationship between engine size and price with respect to fuel type
ggplot(data = auto_data, aes(x = price, y = engine_size)) +
  geom_point(aes(color = fuel_type, shape = fuel_type)) +
  ggtitle('Relationship between engine size and price with respect to fuel type')
```


Relationship between engine size and price with respect to fuel type



```
# creating heatmap since R doesn't have good visualization for heat map
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# # Load the transformed data from converted from R
transformed_data = pd.read_csv("transformed_data.csv")
corr_matrix = transformed_data.corr()
corr_matrix = round(corr_matrix,2)

# Plot the correlation
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", square=True)
plt.title("Feature Correlation Heatmap")
```



```
# hypothesis testing
```

```
# hypothesis test: does high horsepower means higher price?
```

```
cor.test(auto_data$horsepower,auto_data$price)
```

```
##
```

```
## Pearson's product-moment correlation
```

```
##
```

```
## data: auto_data$horsepower and auto_data$price
```

```
## t = 16.152, df = 203, p-value < 2.2e-16
```

```
## alternative hypothesis: true correlation is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 0.6830811 0.8043001
```

```
## sample estimates:
```

```
## cor
```

```
## 0.7499191
```

```
# indicates the good correlation between horsepower and price
```

```
# hypothesis test: does lengthier car has higher price?
```

```
cor.test(auto_data$length,auto_data$price)
```

```
##
## Pearson's product-moment correlation
##
## data: auto_data$length and auto_data$price
## t = 13.454, df = 203, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.6066040 0.7527782
## sample estimates:
## cor
## 0.6865674
```

indicates the good correlation between length and price

does engine size impact horse power
`cor.test(auto_data$engine_size,auto_data$horsepower)`

```
##
## Pearson's product-moment correlation
##
## data: auto_data$engine_size and auto_data$horsepower
## t = 19.695, df = 203, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.7572577 0.8525897
## sample estimates:
## cor
## 0.810216
```

indicates the strong correlation between engine_size and horsepower

does engine size impact compression ratio
`cor.test(auto_data$engine_size,auto_data$compression_ratio)`

```
##
## Pearson's product-moment correlation
##
## data: auto_data$engine_size and auto_data$compression_ratio
## t = 0.41295, df = 203, p-value = 0.6801
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.1084944 0.1653499
## sample estimates:
## cor
## 0.02897136
```

there is no statistically significant evidence that engine size impact compression ratio does accepti

Task 2

Attached is the sales dataset for retail store. Perform data munging on the given dataset.

```
library(readr)
library(reshape2)
library(tidyverse)
library(RColorBrewer)

# Load the data
df <- read_csv("Sales.csv")
```

Basic Data Exploration

```
# Print the shape of the data
dim(df)
```

```
## [1] 1000 17
```

```
head(df)
```

```
## # A tibble: 6 x 17
##   'Invoice ID' Branch City   'Customer type' Gender 'Product line' 'Unit price'
##   <chr>         <chr> <chr>   <chr>          <chr> <chr>          <dbl>
## 1 750-67-8428  A      Yangon Member        Female Health and be~    74.7
## 2 226-31-3081  C      Naypyi~ Normal        Female Electronic ac~    15.3
## 3 631-41-3108  A      Yangon Normal        Male   Home and life~    46.3
## 4 123-19-1176  A      Yangon Member        Male   Health and be~    58.2
## 5 373-73-7910  A      Yangon Normal        Male   Sports and tr~    86.3
## 6 699-14-3026  C      Naypyi~ Normal        Male   Electronic ac~    85.4
## # i 10 more variables: Quantity <dbl>, 'Tax 5%' <dbl>, Total <dbl>, Date <chr>,
## #   Time <time>, Payment <chr>, cogs <dbl>, 'gross margin percentage' <dbl>,
## #   'gross income' <dbl>, Rating <dbl>
```

```
# Check column info of the data
str(df)
```

```
## spc_tbl_ [1,000 x 17] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ Invoice ID      : chr [1:1000] "750-67-8428" "226-31-3081" "631-41-3108" "123-19-1176" ...
## $ Branch         : chr [1:1000] "A" "C" "A" "A" ...
## $ City           : chr [1:1000] "Yangon" "Naypyitaw" "Yangon" "Yangon" ...
## $ Customer type  : chr [1:1000] "Member" "Normal" "Normal" "Member" ...
## $ Gender         : chr [1:1000] "Female" "Female" "Male" "Male" ...
## $ Product line   : chr [1:1000] "Health and beauty" "Electronic accessories" "Home and life" ...
## $ Unit price     : num [1:1000] 74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity       : num [1:1000] 7 5 7 8 7 7 6 10 2 3 ...
## $ Tax 5%         : num [1:1000] 26.14 3.82 16.22 23.29 30.21 ...
## $ Total          : num [1:1000] 549 80.2 340.5 489 634.4 ...
## $ Date           : chr [1:1000] "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" ...
## $ Time           : 'hms' num [1:1000] 13:08:00 10:29:00 13:23:00 20:33:00 ...
## ..- attr(*, "units")= chr "secs"
## $ Payment        : chr [1:1000] "Ewallet" NA "Credit card" "Ewallet" ...
## $ cogs           : num [1:1000] 522.8 76.4 324.3 465.8 604.2 ...
## $ gross margin percentage: num [1:1000] 4.76 4.76 4.76 4.76 4.76 ...
## $ gross income   : num [1:1000] 26.14 3.82 16.22 23.29 30.21 ...
## $ Rating         : num [1:1000] 9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
```

```
## - attr(*, "spec")=
## .. cols(
##   'Invoice ID' = col_character(),
##   Branch = col_character(),
##   City = col_character(),
##   'Customer type' = col_character(),
##   Gender = col_character(),
##   'Product line' = col_character(),
##   'Unit price' = col_double(),
##   Quantity = col_double(),
##   'Tax 5%' = col_double(),
##   Total = col_double(),
##   Date = col_character(),
##   Time = col_time(format = ""),
##   Payment = col_character(),
##   cogs = col_double(),
##   'gross margin percentage' = col_double(),
##   'gross income' = col_double(),
##   Rating = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
spec(df)
```

```
## cols(
##   'Invoice ID' = col_character(),
##   Branch = col_character(),
##   City = col_character(),
##   'Customer type' = col_character(),
##   Gender = col_character(),
##   'Product line' = col_character(),
##   'Unit price' = col_double(),
##   Quantity = col_double(),
##   'Tax 5%' = col_double(),
##   Total = col_double(),
##   Date = col_character(),
##   Time = col_time(format = ""),
##   Payment = col_character(),
##   cogs = col_double(),
##   'gross margin percentage' = col_double(),
##   'gross income' = col_double(),
##   Rating = col_double()
## )
```

Data Munging

```
# Check for missing values
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##                10                0                0
##      Customer type           Gender      Product line
##                0                0                0
```

```
##          Unit price          Quantity          Tax 5%
##              2              16              14
##          Total          Date          Time
##              14              0              1
##          Payment          cogs gross margin percentage
##              3              0              29
##          gross income          Rating
##              0              0
```

Imputing the missing values

```
# imputing missing values for invoice
impute_invoice <- function() {
  x <- sample(c(100:999), 10, replace = T)
  y <- sample(c(10:99), 10, replace = T)
  z <- sample(c(100:999), 10, replace = T)
  a <- paste0(x, "-", y, "-", z)
  return (a)
}

df[is.na(df$'Invoice ID'),1] <- impute_invoice()
head(df)
```

```
## # A tibble: 6 x 17
##   'Invoice ID' Branch City   'Customer type' Gender 'Product line' 'Unit price'
##   <chr>         <chr> <chr>   <chr>          <chr> <chr>          <dbl>
## 1 750-67-8428  A     Yangon Member      Female Health and be~    74.7
## 2 226-31-3081  C     Naypyi~ Normal      Female Electronic ac~    15.3
## 3 631-41-3108  A     Yangon Normal      Male    Home and life~    46.3
## 4 123-19-1176  A     Yangon Member      Male    Health and be~    58.2
## 5 373-73-7910  A     Yangon Normal      Male    Sports and tr~    86.3
## 6 699-14-3026  C     Naypyi~ Normal      Male    Electronic ac~    85.4
## # i 10 more variables: Quantity <dbl>, 'Tax 5%' <dbl>, Total <dbl>, Date <chr>,
## #   Time <time>, Payment <chr>, cogs <dbl>, 'gross margin percentage' <dbl>,
## #   'gross income' <dbl>, Rating <dbl>
```

```
colSums(is.na(df))
```

```
##          Invoice ID          Branch          City
##              0              0              0
##          Customer type          Gender          Product line
##              0              0              0
##          Unit price          Quantity          Tax 5%
##              2              16              14
##          Total          Date          Time
##              14              0              1
##          Payment          cogs gross margin percentage
##              3              0              29
##          gross income          Rating
##              0              0
```

```
# missing value for Invoice is filled.
```



```
# now missing values for unit price
```

```
df[is.na(df$'Unit price'),c('Unit price','Total','Tax 5%','Quantity')]
```

```
## # A tibble: 2 x 4
```

```
##   'Unit price' Total 'Tax 5%' Quantity
##   <dbl> <dbl>    <dbl>    <dbl>
## 1      NA  772.    36.8      10
## 2      NA  125.     5.96      5
```

```
df$'Unit price' <- ifelse(is.na(df$'Unit price'), (df$Total - df$'Tax 5%') / df$Quantity, df$'Unit price')
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##              0              0              0
##   Customer type           Gender       Product line
##              0              0              0
##           Unit price       Quantity       Tax 5%
##              0             16             14
##           Total           Date           Time
##             14              0              1
##           Payment       cogs gross margin percentage
##              3              0              29
##   gross income           Rating
##              0              0
```

```
# missing value for unit price is filled.
```

```
# missing values for Quantity
```

```
head(df[is.na(df$'Quantity'),])
```

```
## # A tibble: 6 x 17
```

```
##   'Invoice ID' Branch City   'Customer type' Gender 'Product line' 'Unit price'
##   <chr>      <chr> <chr>   <chr>          <chr> <chr>          <dbl>
## 1 418-05-0656 B      Mandal~ Normal      Female Fashion acces~    25.6
## 2 804-38-3935 A      Yangon Member      Male  Electronic ac~    93.8
## 3 866-70-2814 B      Mandal~ Normal      Female Electronic ac~    52.8
## 4 101-81-4070 C      Naypyi~ Member      Female Health and be~    62.8
## 5 851-98-3555 B      Mandal~ Normal      Female Health and be~    82.9
## 6 186-71-5196 A      Yangon Member      Female Food and beve~    79.5
## # i 10 more variables: Quantity <dbl>, 'Tax 5%' <dbl>, Total <dbl>, Date <chr>,
## #   Time <time>, Payment <chr>, cogs <dbl>, 'gross margin percentage' <dbl>,
## #   'gross income' <dbl>, Rating <dbl>
```

```
df$'Quantity' <- ifelse(is.na(df$'Quantity'), (df$Total - df$'Tax 5%') %/% df$'Unit price', df$'Quantity')
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##              0              0              0
##   Customer type           Gender       Product line
##              0              0              0
##           Unit price       Quantity       Tax 5%
##              0              0             14
```

```
##           Total           Date           Time
##           14             0             1
##           Payment        cogs gross margin percentage
##           3              0             29
##           gross income    Rating
##           0              0
```

missing value for Quantity is filled.

missing values for Total

```
df[is.na(df$'Total'),c('Unit price','Total','Tax 5%','Quantity')]
```

```
## # A tibble: 14 x 4
##   'Unit price' Total 'Tax 5%' Quantity
##   <dbl> <dbl> <dbl> <dbl>
## 1      88.4   NA    22.1     5
## 2      44.6   NA    11.1     5
## 3      23.1   NA    10.4     9
## 4      66.1   NA    13.2     4
## 5      80.0   NA    20.0     5
## 6      74.3   NA     3.71     1
## 7      24.8   NA     6.19     5
## 8      24.9   NA    11.2     9
## 9      54.4   NA     2.72     1
## 10     48.5   NA    17.0     7
## 11     99.4   NA     9.94     2
## 12     25.4   NA    10.2     8
## 13     21.6   NA     1.08     1
## 14     82.6   NA    41.3    10
```

```
df$'Total' <- ifelse(is.na(df$'Total'), (df$'Unit price' * df$'Quantity') + df$'Tax 5%', df$'Total')
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##           0             0             0
##           Customer type        Gender           Product line
##           0             0             0
##           Unit price           Quantity           Tax 5%
##           0             0             14
##           Total           Date           Time
##           0             0             1
##           Payment        cogs gross margin percentage
##           3              0             29
##           gross income    Rating
##           0              0
```

missing value for Total is filled.

missing values for Tax 5%

```
df[is.na(df$'Tax 5%'),c('Unit price','Total','Tax 5%','Quantity')]
```

```
## # A tibble: 14 x 4
##   'Unit price' Total 'Tax 5%' Quantity
##   <dbl> <dbl>    <dbl>    <dbl>
## 1      93.7 590.      NA        6
## 2      68.9 507.      NA        7
## 3      72.6 457.      NA        6
## 4      89.5 940.      NA       10
## 5      62.1 652.      NA       10
## 6      48.5 153.      NA         3
## 7      87.9 923.      NA       10
## 8      57.1 420.      NA         7
## 9      72.5 609      NA         8
## 10     43.2 363.      NA         8
## 11     15.3 16.1      NA         1
## 12     23.5 49.3      NA         2
## 13     75.8 79.6      NA         1
## 14     73.0 767.      NA       10

df$'Tax 5%' <- ifelse(is.na(df$'Tax 5%'), (df$'Total' - (df$'Unit price') * df$'Quantity'), df$'Total')
colSums(is.na(df))

##           Invoice ID           Branch           City
##           0           0           0
##   Customer type       Gender       Product line
##           0           0           0
##   Unit price       Quantity       Tax 5%
##           0           0           0
##   Total           Date           Time
##           0           0           1
##   Payment       cogs gross margin percentage
##           3           0           29
##   gross income       Rating
##           0           0

# missing value for Tax 5% is filled.

# missing values for Time
df[is.na(df$Time),c('Date', 'Time')]

## # A tibble: 1 x 2
##   Date       Time
##   <chr>    <time>
## 1 2/9/2019    NA

random_time <- as.character(sample(df$Time,1))
df[is.na(df$Time),'Time'] <- random_time
df$Time <- gsub(":00", "",df$Time)

colSums(is.na(df))

##           Invoice ID           Branch           City
##           0           0           0
```

```
##           Customer type           Gender           Product line
##           0                0                0
##           Unit price           Quantity           Tax 5%
##           0                0                0
##           Total                Date                Time
##           0                0                0
##           Payment              cogs gross margin percentage
##           3                0                29
##           gross income          Rating
##           0                0
```

missing value for Time is filled.

missing values for Payment

```
df[is.na(df$Payment),]
```

```
## # A tibble: 3 x 17
```

```
##   'Invoice ID' Branch City   'Customer type' Gender 'Product line' 'Unit price'
##   <chr>         <chr> <chr>   <chr>         <chr> <chr>         <dbl>
## 1 226-31-3081  C     Naypyi~ Normal      Female Electronic ac~      15.3
## 2 145-94-9061  B     Mandal~ Normal      Female Food and beve~      88.4
## 3 841-35-6630  C     Naypyi~ Normal      Female Electronic ac~      75.9
## # i 10 more variables: Quantity <dbl>, 'Tax 5%' <dbl>, Total <dbl>, Date <chr>,
## #   Time <chr>, Payment <chr>, cogs <dbl>, 'gross margin percentage' <dbl>,
## #   'gross income' <dbl>, Rating <dbl>
```

```
unique(df$Payment)
```

```
## [1] "Ewallet"      NA              "Credit card" "Cash"
```

```
most_payment <- names(which.max(table(df$Payment)))
```

```
df[is.na(df$Payment), 'Payment'] <- most_payment
```

```
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##           0                0                0
##           Customer type           Gender           Product line
##           0                0                0
##           Unit price           Quantity           Tax 5%
##           0                0                0
##           Total                Date                Time
##           0                0                0
##           Payment              cogs gross margin percentage
##           0                0                29
##           gross income          Rating
##           0                0
```

missing value for Tax 5% is Payment

missing values for gross margin percentage

```
df[is.na(df$'gross margin percentage'),
  c('Total', 'gross income', 'gross margin percentage')]
```

```
## # A tibble: 29 x 3
##   Total 'gross income' 'gross margin percentage'
##   <dbl>         <dbl>         <dbl>
## 1  23.5           1.12           NA
## 2  177.           8.45           NA
## 3  102.           4.84           NA
## 4  349.          16.6           NA
## 5   85.5          4.07           NA
## 6  336.          16.0           NA
## 7  217.          10.3           NA
## 8  175.           8.33           NA
## 9  335.          16.0           NA
## 10 329.          15.7           NA
## # i 19 more rows
```

```
df[,c('Total', 'gross income', 'gross margin percentage')]
```

```
## # A tibble: 1,000 x 3
##   Total 'gross income' 'gross margin percentage'
##   <dbl>         <dbl>         <dbl>
## 1  549.          26.1           4.76
## 2   80.2          3.82           4.76
## 3  341.          16.2           4.76
## 4  489.          23.3           4.76
## 5  634.          30.2           4.76
## 6  628.          29.9           4.76
## 7  434.          20.7           4.76
## 8  772.          36.8           4.76
## 9   76.1          3.63           4.76
## 10 173.           8.23           4.76
## # i 990 more rows
```

```
df[is.na(df$'gross margin percentage'),'gross margin percentage'] <- unique(df$'gross margin percentage')
colSums(is.na(df))
```

```
##           Invoice ID           Branch           City
##           0           0           0
##   Customer type           Gender           Product line
##           0           0           0
##   Unit price           Quantity           Tax 5%
##           0           0           0
##           Total           Date           Time
##           0           0           0
##           Payment           cogs gross margin percentage
##           0           0           0
##           gross income           Rating
##           0           0
```

```
# missing value for gross margin percentage.
```


```
# Convert the date column to datetime format
df$Date <- as.Date(mdy(df$Date))
```

Basic EDA on Sales Dataset

```
# Plot the distribution of different branches
branch_counts <- table(df$Branch)

colors <- brewer.pal(length(branch_counts), "Paired")
explode <- rep(0, length(branch_counts))
explode[which.max(branch_counts)] <- 0.1

ggplot(data.frame(Branch = names(branch_counts), Count = as.numeric(branch_counts)),
  aes(x = "", y = Count, fill = Branch)) +
  geom_bar(stat = "identity", width = 1, color = "white") +
  coord_polar("y", start = 0) +
  geom_text(aes(label = paste0(round(Count / sum(branch_counts) * 100, 1), "%")),
    position = position_stack(vjust = 0.5)) +
  scale_fill_manual(values = colors) +
  theme_void() +
  theme(legend.position = "right",
    legend.title = element_blank(),
    plot.title = element_text(hjust = 0.5, size = 16, face = "bold"),
    plot.subtitle = element_text(hjust = 0.5, size = 12)) +
  labs(title = "Branch Distribution")
```




```
# Plot the distribution of different cities
ggplot(df, aes(x = City, fill = City)) +
  geom_bar(width = 0.6, color = "white") +
```

```

scale_fill_manual(values = brewer.pal(length(unique(df$City)), "Paired")) +
labs(title = "Distribution of Different Cities",
     x = "City",
     y = "Count") +
geom_text(stat = "count", aes(label = after_stat(count)), vjust = -0.5, size = 3.5) +
ylim(0, 400) +
theme_minimal() +
theme(plot.title = element_text(size = 16, face = "bold", hjust = 0.5),
      axis.text.x = element_text(angle = 45, hjust = 1, size = 8),
      axis.title = element_text(size = 12),
      legend.position = "right",
      legend.title = element_blank())

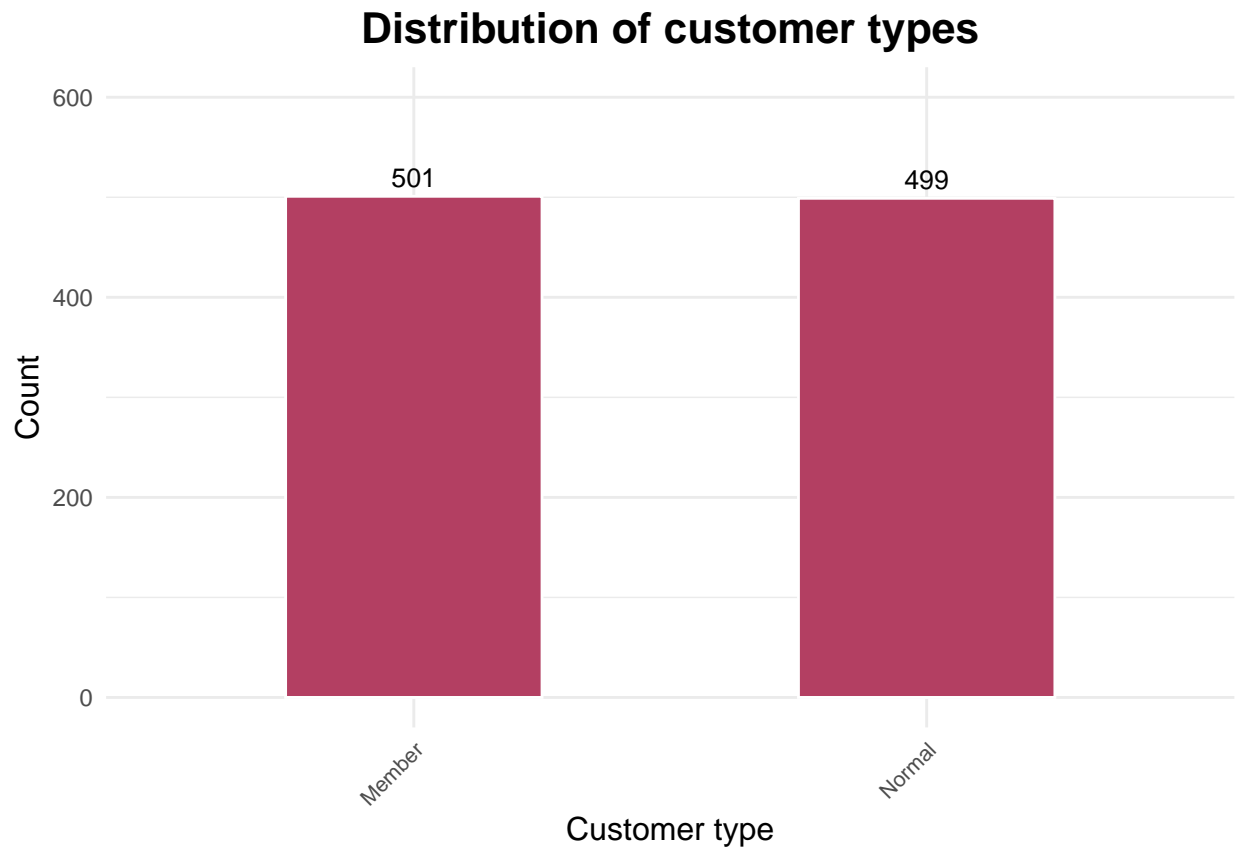
```



```


# Plot the distribution of customer types
ggplot(df, aes(x = df$'Customer type')) +
geom_bar(width = 0.5, color = "white", fill = "#B33F62") +
labs(title = "Distribution of customer types", x = "Customer type", y = "Count") +
geom_text(stat = "count", aes(label = comma(..count..)), vjust = -0.5, size = 3.5) +
ylim(0, 600) +
theme_minimal() +
theme(plot.title = element_text(size = 16, face = "bold", hjust = 0.5),
      axis.text.x = element_text(angle = 45, hjust = 1, size = 8),
      axis.title = element_text(size = 12),
      legend.position = "right",
      legend.title = element_blank())

```



```
# Plot the distribution of genders
ggplot(data.frame(Gender = names(table(df$Gender)), Count = as.numeric(table(df$Gender))), aes(x = "", y = Count)) +
  geom_bar(stat = "identity", width = 1, color = "white") +
  coord_polar("y", start = 0) +
  geom_text(aes(label = paste0(round(Count / sum(branch_counts) * 100, 1), "%")),
            position = position_stack(vjust = 0.5)) +
  scale_fill_manual(values = colors) +
  theme_void() +
  theme(legend.position = "right",
        legend.title = element_blank(),
        plot.title = element_text(hjust = 0.5, size = 16, face = "bold"),
        plot.subtitle = element_text(hjust = 0.5, size = 12)) +
  labs(title = "Gender Distribution")
```



Gender Distribution



```
# Product Line
table(df$'Product line')


##
## Electronic accessories    Fashion accessories    Food and beverages
##           170              178              174
##   Health and beauty    Home and lifestyle    Sports and travel
##           152              160              166

ggplot(df, aes(x = df$'Product line', fill = df$'Product line')) +
  geom_bar(width = 0.6, color = "white") +
  scale_fill_manual(values = brewer.pal(length(unique(df$'Product line')), "Paired")) +
  labs(title = "Distribution of Different Cities",
       x = "City",
       y = "Count") +
  geom_text(stat = "count", aes(label = after_stat(count)), vjust = -0.5, size = 3.5) +
  ylim(0, 200) +
  theme_minimal() +
  theme(plot.title = element_text(size = 16, face = "bold", hjust = 0.5),
        axis.text.x = element_text(angle = 45, hjust = 1, size = 8),
        axis.title = element_text(size = 12),
        legend.position = "right",
        legend.title = element_blank())
```




```
# Plot the distribution of unit prices  
plot(density(df$'Unit price',na.rm = T), main = 'Distribution of Unit Price')
```

Distribution of Unit Price



```
# plot scatterplot between total and rating with respect to Branch  
ggplot(data=df, aes(x= Total, y= Rating))+  
  geom_point(aes(color=Branch, shape=Branch))
```



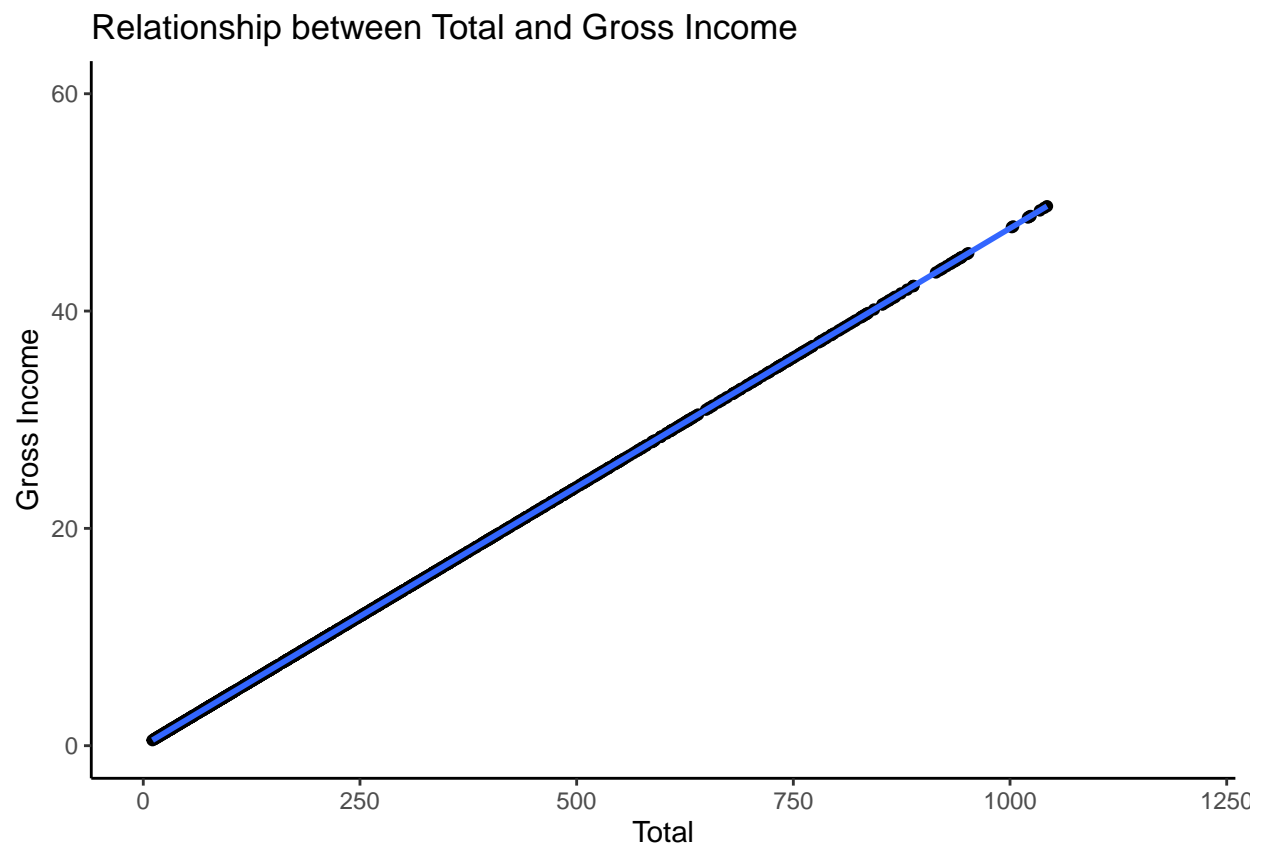
```
# Hypothesis testing: Does gross income increase when total sales increase?
cor.test(df$Total, df$'gross income')
```

```
##
## Pearson's product-moment correlation
##
## data: df$Total and df$"gross income"
## t = Inf, df = 998, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  1 1
## sample estimates:
## cor
##  1
```

```
# it is indeed true, and is statistically significant
```


```
# Plot the relationship between total sales and gross income
ggplot(df, aes(x = Total, y = df$`gross income`)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(
    title = "Relationship between Total and Gross Income",
    x = "Total",
    y = "Gross Income"
```

```
) +  
theme_classic() +  
xlim(0,1200) + ylim(0,60)
```



```
# Plot the relationship between customer type and total sales  
ggplot(df, aes(x = Total, color = df$`Customer type`)) +  
  geom_density() +  
  labs(title = "Relationship between customer type and total sales", x = "Total sales", y = "Density")
```

Relationship between customer type and total sales



Which branch sales has the highest ratings?

```
df %>%
  group_by(Branch) %>%
  summarise(mean_rating = mean(Rating)) %>%
  arrange(desc(mean_rating))
```

```
## # A tibble: 3 x 2
##   Branch mean_rating
##   <chr>         <dbl>
## 1 C             7.07
## 2 A             7.03
## 3 B             6.82
```

which product is used by different group of Gender

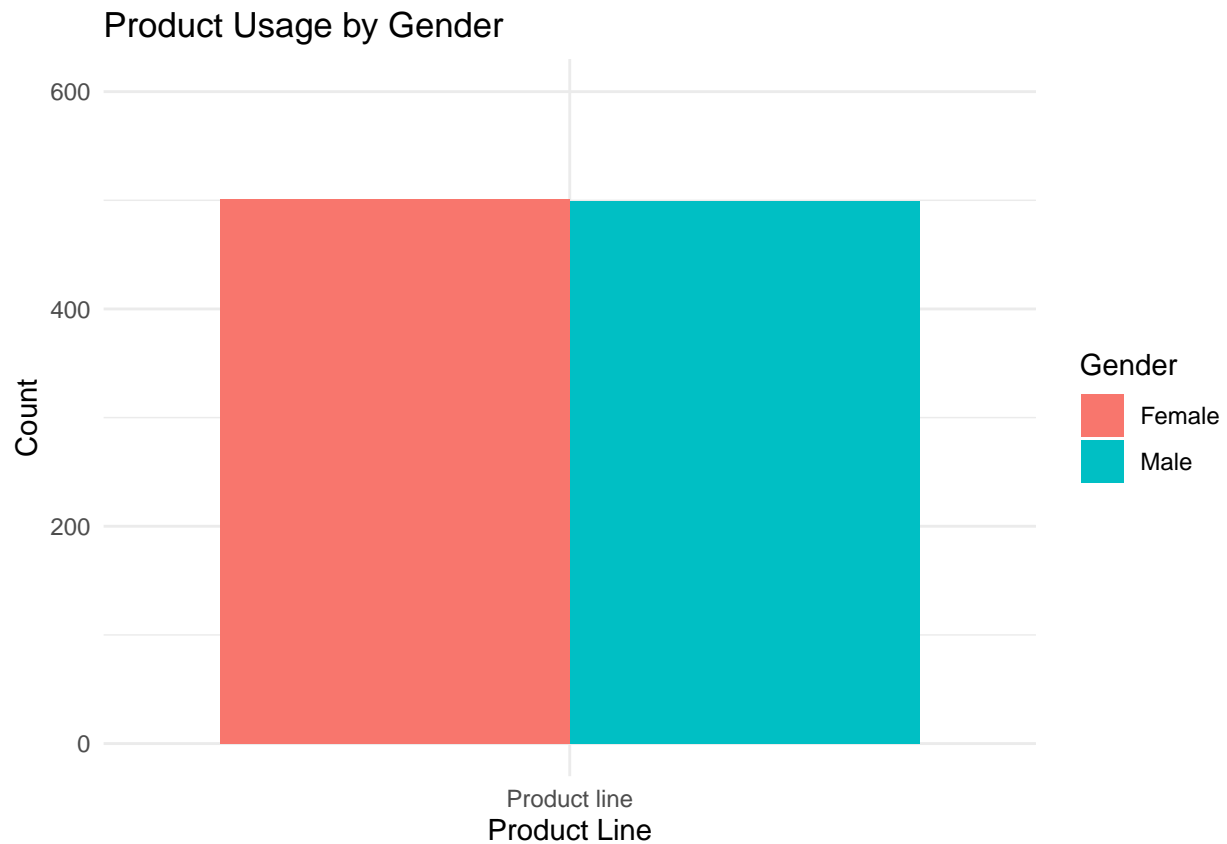
```
product_usage <- df %>%
  group_by(Gender, 'Product line') %>%
  summarize(Count = n()) %>%
  arrange(desc(Count))
```

```
ggplot(product_usage, aes(x = 'Product line', y = Count,
  fill = Gender)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(
    title = "Product Usage by Gender",
    x = "Product Line",
```

```

  y = "Count",
  fill = "Gender"
) +
theme_minimal() +
theme(
  legend.position = "right"
)+ ylim(0,600)

```



```

product_usage_city <- df %>%
  group_by(City, 'Product line') %>%
  summarize(Count = n()) %>%
  arrange(desc(Count))

```

'summarise()' has grouped output by 'City'. You can override using the
'.groups' argument.

```

names(product_usage_city) <- c('City', 'Product', 'Count')
product_usage_city

```

```

## # A tibble: 3 x 3
## # Groups:   City [3]
##   City      Product      Count
##   <chr>    <chr>      <int>
## 1 Yangon   Product line    340
## 2 Mandalay Product line    332
## 3 Naypyitaw Product line    328

```

Task 3

Refer the Opinion published on Himalayan times on Dec 19, 2023 and perform a text preprocessing and generate word cloud.

```
#
import requests
from bs4 import BeautifulSoup
url = 'https://thehimalayantimes.com/opinion/navigating-nepals-digital-frontier-\
understanding-cybersecurity-in-the-digital-age-ensuring-data-safety-and-the-role-of-ai'

x = requests.get(url)
soup = BeautifulSoup(x.content, 'html.parser')
post_content = soup.find('div', {'class': 'post-content'})
paragraphs = post_content.find_all('p')
final_list = ''
for paragraph in range(0, len(paragraphs)-2):
    final_list += (paragraphs[paragraph].text)

with open('himalayan_times.txt','w') as file:
    file.write(final_list)

## 5513

library(tm)
library(Rgraphviz)
library(wordcloud)

text_document <- readLines('himalayan_times.txt')
corpus <- Corpus(VectorSource(text_document))
```

Text Preprocessing:

- Remove Punctutaion
- Remove Stop Words
- Stemming
- Convert to Lower
- Remove any Numbers
- Any customer remove words

```
my_stopwords <- c("can", "may", "used")
corpus <- tm_map(corpus, removeWords, my_stopwords)
my_tdm <- TermDocumentMatrix(
  corpus,
  control =
    list(
      removePunctuation = TRUE,
      stopwords = TRUE,
      tolower = TRUE,
      stemming = FALSE,
      removeNumbers = TRUE,
      bounds = list(global = c(1, Inf)),
      wordLenghts = c(1, Inf),
      removeWords = (c("can", "may", "used")))
)
```



```

# find the frequent terms in the corpus
frequent_terms <- findFreqTerms(my_tdm)
head(frequent_terms,20)

## [1] "ability"      "access"      "accessed"    "achieve"
## [5] "additionally" "adoption"    "advances"    "advent"
## [9] "ais"          "aithe"       "aligned"     "alikebuilding"
## [13] "allocated"    "allocating"  "allowing"    "along"
## [17] "already"      "also"        "always"      "amounts"

mat <- as.matrix(my_tdm)
freq <- mat %>% rowSums() %>% sort(decreasing = T)


df <- my_tdm %>%
  as.matrix() %>%
  rowSums() %>%
  sort(decreasing = TRUE) %>%
  head(10) %>%
  enframe(name = "word", value = "counts")
head(df)

## # A tibble: 6 x 2
##   word      counts
##   <chr>      <dbl>
## 1 digital      18
## 2 security     11
## 3 cyber         6
## 4 smart         6
## 5 data          5
## 6 framework     5

# top 10 words and counts using bargraph
library(ggplot2)
ggplot(df, aes(word, counts)) +
  geom_bar(stat = "identity", fill = "#932421") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs(title = "Top 10 words by counts.") +
  geom_text(aes(label = counts), vjust = -0.5)

```

Top 10 words by counts.



```
# plot word cloud
wordcloud(
  words = names(freq),
  freq = freq,
  random.order = FALSE,
  colors = brewer.pal(8, "Dark2"),
  scale = c(4, 0.5),
  random.color = TRUE,
)
grid.text("Himalayan Times Word Cloud", x = 0.5, y = 0.9, gp = gpar(fontsize = 18, fontface = "bold"))
```

Himalayan Times Word Cloud

