

# Statistical Computing with R

## Masters in Data Science 503 (S8&9)

### Third Batch, SMS, TU, 2024

Shital Bhandary

Associate Professor

Statistics/Bio-statistics, Demography and Public Health Informatics

Patan Academy of Health Sciences, Lalitpur, Nepal

Faculty, Data Analysis and Decision Modeling, MBA, Pokhara University, Nepal

Faculty, FAIMER Fellowship in Health Professions Education, India/USA.

# Review Preview (Unit 2, Part 2 & 3)

- Data wrangling
- Data munching
- Tidy data
- “dplyr” package and its use for data manipulation
- Reading database in R
- Big data in R
- Text Mining

# Data wrangling (Course book Chapter 9-16)

R for Data Science, <https://r4ds.had.co.nz/index.html>

- Data wrangling is the art of getting your data into R in a useful form for visualization and modelling.
- Data wrangling is very important: without it you can't work with your own data! There are three main parts to data wrangling:
  - Import
  - Tidy
  - Transform

# Import data in R:

- We have already covered this in the previous classes
- More here: <https://r4ds.had.co.nz/data-import.html>
- Reach this chapter well as there are some important import functions that are part of this course and may not have discussed so far
- We will discuss about reading “database” in the second part of this class

# Tidy data in R

- Tidy data is a consistent way to organize your data in R.
- Getting your/our data into this format requires some upfront work, but that work pays off in the long term.
- Once you/we have tidy data and the tidy tools provided by packages in the **tidyverse**, you will spend much less time **munging/cleaning data** from one representation to another, allowing you to spend more time on the analytic questions at hand.
- Tidy data in tidyverse packages are stored as “tibble”

Let us see what is “tibble” first:

<https://r4ds.had.co.nz/tibbles.html>

- The variant of the data frame used by “tidiverse” is called: **tibble**.
- Tibbles are data frames, but they tweak some older behaviours to make life a little easier.
- R is an old language, and some things that were useful 10 or 20 years ago now get in your way.
- It's difficult to change base R without breaking existing code, so **most innovation occurs in packages**.
- The **tibble** package provides opinionated data frames that make working in the **tidyverse** a little easier.
- It's particularly **useful for large datasets** because it only prints the first few rows.

# Note:

- All the functions of “tidyverse” package works fast with tibble so it will be wise to say that data frame/s should be converted to tibble before using functions of “tidyverse” package
- However, most of the packages of the “tidyverse” super package works well with the data frame too!
- There are two main differences in the usage of a data frame vs a tibble: printing, and subsetting. <https://posit.co/blog/tibble-1-0-0/>

# There are three interrelated rules which make a dataset tidy:

- Each variable must have its own column.
- Each observation must have its own row.
- Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	2745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

variables

country	year	cases	population
Afghanistan	1999	2745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

observations

country	year	cases	population
Afghanistan	99	745	19987071
Afghanistan	00	666	20595360
Brazil	99	37737	172006362
Brazil	00	80488	174504898
China	99	212258	1272915272
China	00	216766	128042583

values



# Example: Which one is “tidy”? Why?

#Creating tibble:

```
table1 <- tibble(  
  country = c("Afghanistan", "Afghanistan", "Brazil",  
    "Brazil", "China", "China"),  
  year = c(1999, 2000, 1999, 2000, 1999, 2000),  
  cases = c(745, 2666, 37737, 80488, 212258, 213766),  
  population = c(19987071, 20595360, 172006362,  
    174504898, 1272915272, 1280428583)  
)
```

# Data frame to tibble: `as_tibble(data_frame)`

# Tibble to data frame: `as.data.frame(tibble_data)`

- table1
- A tibble: 6 x 4
- |     | year  | country     | cases  | population |
|-----|-------|-------------|--------|------------|
|     | <dbl> | <chr>       | <dbl>  | <dbl>      |
| • 1 | 1999  | Afghanistan | 745    | 19987071   |
| • 2 | 2000  | Afghanistan | 2666   | 20595360   |
| • 3 | 1999  | Brazil      | 37737  | 172006362  |
| • 4 | 2000  | Brazil      | 80488  | 174504898  |
| • 5 | 1999  | China       | 212258 | 1272915272 |
| • 6 | 2000  | China       | 213766 | 1280428583 |

- **dbl = duple instead of number in “tibble”!**

# Example: Which one is “tidy”? Why?

- table2
- #> # A tibble: 12 × 4
- #> country year type count
- #> <chr> <int> <chr> <int>
- #> 1 Afghanistan 1999 cases 745
- #> 2 Afghanistan 1999 population 19987071
- #> 3 Afghanistan 2000 cases 2666
- #> 4 Afghanistan 2000 population 20595360
- #> 5 Brazil 1999 cases 37737
- #> 6 Brazil 1999 population 172006362
- #> # ... with 6 more rows

- table3
- #> # A tibble: 6 × 3
- #> country year rate
- #> \* <chr> <int> <chr>
- #> 1 Afghanistan 1999 745/19987071
- #> 2 Afghanistan 2000 2666/20595360
- #> 3 Brazil 1999 37737/172006362
- #> 4 Brazil 2000 80488/174504898
- #> 5 China 1999 212258/1272915272
- #> 6 China 2000 213766/1280428583

# Example: Which one is “tidy”? Why?

## # Spread across two tibbles

### table4a # cases

- #> # A tibble: 3 × 3
- #> country       `1999`       `2000`
- #> \* <chr>       <int>       <int>
- #> 1 Afghanistan     745       2666
- #> 2 Brazil         37737      80488
- #> 3 China          212258     213766

### table4b # population

- #> # A tibble: 3 × 3
- #> country       `1999`       `2000`
- #> \* <chr>       <int>       <int>
- #> 1 Afghanistan 19987071   20595360
- #> 2 Brazil       172006362   174504898
- #> 3 China       1272915272 1280428583

# Why ensure that your data is tidy? There are two main advantages:

- There's a general advantage to picking one consistent way of storing data. If you have a consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
- There's a specific advantage to placing variables in columns because it allows R's vectorized nature to shine.
- **dplyr**, **ggplot2**, and all the other packages in the **tidyverse** are designed to work with tidy data.

# Tidy data: Pivoting – Longer to wider

## (To do standard statistical analysis)

- table2
- #> # A tibble: 12 × 4
- #> country year type count
- #> <chr> <int> <chr> <int>
- #> 1 Afghanistan 1999 cases 745
- #> 2 Afghanistan 1999 population 19987071
- #> 3 Afghanistan 2000 cases 2666
- #> 4 Afghanistan 2000 population 20595360
- #> 5 Brazil 1999 cases 37737
- #> 6 Brazil 1999 population 172006362
- #> # ... with 6 more rows

```
table2 %>%  
  pivot_wider(names_from = type, values_from =  
  count)
```

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

  

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table2

Figure 12.3: Pivoting table2 into a wider, tidy form.

# Tidy data: Pivoting – Wider to Longer (To do Variance components analysis)

- table4a
- #> # A tibble: 3 × 3
- #> country `1999` `2000`
- #> \* <chr> <int> <int>
- #> 1 Afghanistan 745 2666
- #> 2 Brazil 37737 80488
- #> 3 China 212258 213766

```
table4a %>%  
  pivot_longer(c(`1999`, `2000`), names_to =  
    "year", values_to = "cases")
```

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

Figure 12.2: Pivoting table4 into a longer, tidy form.

# Tidy data: Separate

- table3
- #> # A tibble: 6 × 3
- #> country year rate
- #> \* <chr> <int> <chr>
- #> 1 Afghanistan 1999 745/19987071
- #> 2 Afghanistan 2000 2666/20595360
- #> 3 Brazil 1999 37737/172006362
- #> 4 Brazil 2000 80488/174504898
- #> 5 China 1999 212258/1272915272
- #> 6 China 2000 213766/1280428583

```
table3 %>%  
  separate(rate, into = c("cases", "population"))  
OR  
table3 %>%  
  separate(rate, into = c("cases", "population"), sep = "/")
```

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

table3

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Figure 12.4: Separating table3 makes it tidy

# Tidy data: Unite

- **unite()** is the inverse of **separate()**: it combines multiple columns into a single column.
- You'll need it much less frequently than **separate()**, but it's still a useful tool to have in your back pocket.

```
table5 %>%  
  unite(new, century, year)  
OR  
table5 %>%  
  unite(new, century, year, sep = "")
```

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	0	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	0	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	0	213766 / 1280428583

table6

Figure 12.5: Uniting table5 makes it tidy



# Missing values

- Changing the representation of a dataset brings up an important subtlety of missing values.
- Surprisingly, a value can be missing in one of two possible ways:
  - **Explicitly**, i.e. flagged with NA.
  - **Implicitly**, i.e. simply not present in the data.

# Missing values: Example

#Create a tibble with missing values:

```
stocks <- tibble(  
  year = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),  
  qtr  = c( 1,  2,  3,  4,  2,  3,  4),  
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)  
)
```

# Missing values: Example

There are two missing values in this dataset:

- The return for the fourth quarter of 2015 is **explicitly missing**, because the cell where its value should be instead contains NA.
- The return for the first quarter of 2016 is **implicitly missing**, because it simply does not appear in the dataset.

# Missing values: Example

- `stocks %>%`
- `pivot_wider(names_from = year, values_from = return)`
- `#> # A tibble: 4 × 3`
- `#>`

	qtr	`2015`	`2016`
<code>#&gt;</code>	<code>&lt;dbl&gt;</code>	<code>&lt;dbl&gt;</code>	<code>&lt;dbl&gt;</code>
<code>#&gt; 1</code>	1	1.88	NA
<code>#&gt; 2</code>	2	0.59	0.92
<code>#&gt; 3</code>	3	0.35	0.17
<code>#&gt; 4</code>	4	NA	2.66
- `#>`
- `#>`
- `#> 1`
- `#> 2`
- `#> 3`
- `#> 4`

# Missing values: What will happen now?

- `stocks %>%`

```
  pivot_wider(names_from = year, values_from = return) %>%
```

```
  pivot_longer(
```

```
    cols = c(`2015`, `2016`),
```

```
    names_to = "year",
```

```
    values_to = "return",
```

```
    values_drop_na = TRUE
```

```
)
```

# Missing values:

## We can use “complete” command!

- `stocks %>%`
- `complete(year, qtr)`
- `#> # A tibble: 8 × 3`
- `#> year qtr return`
- `#> <dbl> <dbl> <dbl>`
- `#> 1 2015 1 1.88`
- `#> 2 2015 2 0.59`
- `#> 3 2015 3 0.35`
- `#> 4 2015 4 NA`
- `#> 5 2016 1 NA`
- `#> 6 2016 2 0.92`
- `#> # ... with 2 more rows`

# Missing values: Another example (tibble by row or tribble!)

- `treatment <- tribble(  
 ~ person, ~ treatment, ~response,  
 "Derrick Whitmore", 1, 7,  
 NA, 2, 10,  
 NA, 3, 9,  
 "Katherine Burke", 1, 4  
)`
  - `treatment`

# Missing values: fill() for another example

- treatment %>%
- fill(person) **# “tidyr” package is required here!**
- #> # A tibble: 4 × 3
- #> person treatment response
- #> <chr> <dbl> <dbl>
- #> 1 Derrick Whitmore 1 7
- #> 2 Derrick Whitmore 2 10
- #> 3 Derrick Whitmore 3 9
- #> 4 Katherine Burke 1 4



Question/Queries so far?

# Transform/manipulate data with “dplyr”

- To learn five key “dplyr” package functions that allow you to solve the vast majority of your data manipulation challenges:
  - Pick observations by their values (**filter()**).
  - Reorder the rows (**arrange()**).
  - Pick variables by their names (**select()**).
  - Create new variables with functions of existing variables (**mutate()**).
  - Collapse many values down to a single summary (**summarise()**).
- These can all be used in conjunction with **group\_by()** which changes the scope of each function from operating on the entire dataset to operating on it group-by-group.

# Data manipulation with “dplyr”

- These six functions provide the verbs for a language of data manipulation.
- All verbs work similarly:
  - The first argument is a data frame.
  - The subsequent arguments describe what to do with the data frame, using the variable names (without quotes).
  - The result is a new data frame.
- Together these properties make it easy to chain together multiple simple steps to achieve a complex result.

# Let's use them with nycflighst13 data

- library(dplyr)
- library(nycflights13)
- flights
- #> # A tibble: 336,776 × 19
- #> year month day dep\_time sched\_dep...<sup>1</sup> dep\_d...<sup>2</sup> arr\_t...<sup>3</sup> sched...<sup>4</sup> arr\_d...<sup>5</sup> carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 336,770 more rows, 9 more variables

# Filter: What will happen?

- **filter(flights, month == 1, day == 1)**
- #> # **A tibble: 842 × 19**
- #> year month day dep\_time sched\_dep...<sup>1</sup> dep\_d...<sup>2</sup> arr\_t...<sup>3</sup> sched...<sup>4</sup>  
arr\_d...<sup>5</sup> carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 836 more rows, 9 more variables

# Are these better?

- `jan1 <- filter(flights, month == 1, day == 1)`
- `(jan1 <- filter(flights, month == 1, day == 1))`
  
- `dec25 <- filter(flights, month == 12, day == 25)`
- `(dec25 <- filter(flights, month == 12, day == 25))`
  
- `filter(flights, month = 1)`      #Why error?
- `filter(flights, month == 1)`      #Works now? Why?

# More with filter:

- `filter(flights, month == 11 | month == 12)` #What?
- `filter(flights, month == 11 | 12)` #Works?
- `nov_dec <- filter(flights, month %in% c(11, 12))` #Works?
- De Morgan's Law:
- `filter(flights, !(arr_delay > 120 | dep_delay > 120))` #Works?
- `filter(flights, arr_delay <= 120, dep_delay <= 120)` #Works?

# Arrange: Example

- `arrange(flights, year, month, day)`
- `#> # A tibble: 336,776 × 19`
- `#> year month day dep_time sched_dep...1 dep_d...2 arr_t...3 sched...4 arr_d...5  
carrier`
- `#> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>`
- `#> 1 2013 1 1 517 515 2 830 819 11 UA`
- `#> 2 2013 1 1 533 529 4 850 830 20 UA`
- `#> 3 2013 1 1 542 540 2 923 850 33 AA`
- `#> 4 2013 1 1 544 545 -1 1004 1022 -18 B6`
- `#> 5 2013 1 1 554 600 -6 812 837 -25 DL`
- `#> 6 2013 1 1 554 558 -4 740 728 12 UA`
- `#> # ... with 336,770 more rows, 9 more variables`



# What will happen now?

- Arrange will sort the data in ascending order
- `arrange(flights, desc(dep_delay))`
- Use `desc()` to re-order by a column in descending order
- Missing values are always sorted at the end

# Select: Example

- # Select columns by name
- `select(flights, year, month, day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

- # Select all columns between year and day (inclusive)
- `select(flights, year:day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

# Select: “except” example

- # Select all columns except those from year to day (inclusive)
- `select(flights, -(year:day))`
- #> # A tibble: 336,776 × 16
- #> dep\_time sched...<sup>1</sup> dep\_d...<sup>2</sup> arr\_t...<sup>3</sup> sched...<sup>4</sup> arr\_d...<sup>5</sup> carrier flight tailnum origin
- #> <int> <int> <dbl> <int> <int> <dbl> <chr> <int> <chr> <chr>
- #> 1 517 515 2 830 819 11 UA 1545 N14228 EWR
- #> 2 533 529 4 850 830 20 UA 1714 N24211 LGA
- #> 3 542 540 2 923 850 33 AA 1141 N619AA JFK
- #> 4 544 545 -1 1004 1022 -18 B6 725 N804JB JFK
- #> 5 554 600 -6 812 837 -25 DL 461 N668DN LGA
- #> 6 554 558 -4 740 728 12 UA 1696 N39463 EWR
- #> # ... with 336,770 more rows, 6 more variables

# Select: More

- There are a number of helper functions you can use within `select()`:
- **`starts_with("abc")`**: matches names that begin with “abc”.
- **`ends_with("xyz")`**: matches names that end with “xyz”.
- **`contains("ijk")`**: matches names that contain “ijk”.
- **`matches("(.)\\1")`**: selects variables that match a **regular expression**.
- This one matches any variables that contain repeated characters.
- **`num_range("x", 1:3)`**: matches x1, x2 and x3.
- See `?select` for more details.

More on regular expression are available here: <https://cran.r-project.org/web/packages/stringr/vignettes/regular-expressions.html>

# Note:

- `select()` can be used to rename variables, **but it's rarely useful because it drops all of the variables not explicitly mentioned.**
- Instead, **use `rename()`**, which is a variant of `select()` that keeps all the variables that aren't explicitly mentioned
- `rename(flights, tail_num = tailnum)`
- Another option is to use `select()` in conjunction with the `everything()` helper.
- This is useful if you have a handful of variables you'd like to move to the start of the data frame.
- `select(flights, time_hour, air_time, everything())`

# Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

```
#Adding variables in flights_sml:
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  speed = distance / air_time * 60
)
```

# Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

#Adding one more variable:

```
mutate(flights_sml,  
      gain = dep_delay - arr_delay,  
      hours = air_time / 60,  
      gain_per_hour = gain / hours  
)
```

#Note that you/we can refer to columns that you've just created

# Transmute and other useful creation functions

More@ <https://r4ds.had.co.nz/transform.html>

- If you only want to keep the new variables, use `transmute()`
  - Arithmetic operators: `+`, `-`, `*`, `/`, `^`
  - Modular arithmetic: `%/%` (integer division) and `%%` (remainder)
  - Use: Compute hour and minute from `dep_time` with:
    - `transmute(flights, dep_time, hour = dep_time %/% 100, minute = dep_time %% 100)`
- ```
transmute(flights,  
  gain = dep_delay - arr_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```



# Summarise: Works best for group summaries

- `summarise(flights, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 1 × 1`
- `#> delay`
- `#> <dbl>`
- `#> 1 12.6`
- `by_day <- group_by(flights, year, month, day)`
- `summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 365 × 4`
- `#> # Groups: year, month [12]`
- `#> year month day delay`
- `#> <int> <int> <int> <dbl>`
- `#> 1 2013 1 1 11.5`
- `#> 2 2013 1 2 13.9`
- `#> 3 2013 1 3 11.0`
- `#> 4 2013 1 4 8.95`
- `#> 5 2013 1 5 5.73`
- `#> 6 2013 1 6 7.15`
- `#> # ... with 359 more rows`

# Multiple operations: pipes

```
#What will happen?
delays <- flights %>%
  group_by(dest) %>%
  summarise(
    count = n(),
    dist = mean(distance, na.rm =
TRUE),
    delay = mean(arr_delay, na.rm =
TRUE)
  ) %>%
  filter(count > 20, dest != "HNL")
```

```
#What will happen?
flights %>%
  group_by(year, month, day) %>%
  summarise(mean =
mean(dep_delay))

#And now?
flights %>%
  group_by(year, month, day) %>%
  summarise(mean =
mean(dep_delay, na.rm = TRUE))
```

# How to remove cancelled flights?

## And, get summaries by groups!

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay),  
         !is.na(arr_delay))  
  
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarise(mean =  
    mean(dep_delay))
```

- #> # A tibble: 365 × 4
- #> # Groups: year, month [12]
- #> year month day mean
- #> <int> <int> <int> <dbl>
- #> 1 2013 1 1 11.4
- #> 2 2013 1 2 13.7
- #> 3 2013 1 3 10.9
- #> 4 2013 1 4 8.97
- #> 5 2013 1 5 5.73
- #> 6 2013 1 6 7.15
- #> # ... with 359 more rows

# Counts: Example

- Whenever you do any aggregation, it's always a good idea to include either a count (`n()`), or a count of non-missing values (`sum(!is.na(x))`).
- That way you can check that you're not drawing conclusions based on very small amounts of data.

# What happens now?

```
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarise(  
    delay = mean(arr_delay)  
  )  
  
hist(delays$delay)
```

# What happens now?

```
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarise(  
    delay = mean(arr_delay,  
na.rm = TRUE),  
    n = n()  
  )
```

# Plots

# Can you interpret them?

```
hist(delays$n)
```

```
hist(delays$delay)
```

```
plot(delays$n, delays$delay)
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# When do the first and last flights leave each day?

```
not_cancelled %>%
```

```
  group_by(year, month, day) %>%
```

```
  summarise(
```

```
    first = min(dep_time),
```

```
    last = max(dep_time)
```

```
  )
```

- # Why is distance to some destinations more variable than to others?

```
not_cancelled %>%
```

```
  group_by(dest) %>%
```

```
  summarise(distance_sd =
```

```
    sd(distance)) %>%
```

```
  arrange(desc(distance_sd))
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# Which destinations have the most carriers?

```
not_cancelled %>%  
  group_by(dest) %>%  
  summarise(carriers =  
n_distinct(carrier)) %>%  
  arrange(desc(carriers))
```

- # How many flights left before 5am? (these usually indicate delayed flights from the previous day)

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarise(n_early =  
sum(dep_time < 500))
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# What proportion of flights are  
delayed by more than an hour?

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarise(hour_prop =  
    mean(arr_delay > 60))
```

#Find all groups bigger than a  
threshold:

```
popular_dests <- flights %>%  
  group_by(dest) %>%  
  filter(n() > 365)  
popular_dests
```



# Popular destination: head and tail

(Are these results VALID?)

- `head(popular_dests$dest)`
- `[1] "IAH" "IAH" "MIA" "BQN"`  
`"ATL" "ORD"`

- IAH = Texas
- MIA = Miami
- BQN = Puerto Rico
- ATL = Atlanta
- ORD = Chicago

- `tail(popular_dests$dest)`
- `[1] "BNA" "DCA" "SYR" "BNA"`  
`"CLE" "RDU"`

- BNA = Nashville
- DCA = Washington (Reagan Nat.)
- SYR = New York (Syracuse)
- CLE = Cleveland
- RDU = North Carolina

# Bonus: dplyr “slice” function with examples

<https://dplyr.tidyverse.org/reference/slice.html>

#What will happen?

`flights %>% slice(1L)`

`flights %>% slice(n())`

`flights %>% slice(5:n())`

`slice(flights,-(1:4))`

- `flights %>% slice_sample(n=5)`
- `flights %>% slice_sample(n=5, replace = TRUE)`
- **set seed(123)**
- `train_data <- flights %>% slice_sample(prop=0.8)`
- `train_data`
- `test_data <- flights %>% slice_sample(prop=0.2)`
- `test_data`

Question/Queries so far?

# Readings for the Next class

- R and Relational database: Chapter 13, R for Data Science, First Edition <https://r4ds.had.co.nz/relational-data.html>
- R for Data Science, 2<sup>nd</sup> Edition, Chapter 22: Databases <https://r4ds.hadley.nz/databases>
- R and Big Data: <https://rviews.rstudio.com/2019/07/17/3-big-data-strategies-for-r/>

# Reading continued ...

## Big Data in R:

- [https://www.columbia.edu/~sjm2186/EPIC\\_R/EPIC\\_R\\_BigData.pdf](https://www.columbia.edu/~sjm2186/EPIC_R/EPIC_R_BigData.pdf)
- Option I: Make the data smaller
  - Issue SQL query directly from R to database
- Option II: Get a bigger computer
- Option III: Use data.table rather than data.frame
- Option IV: Buffer the data set on disk
- Option V: Split it up

Question/Queries?

# Thank you!

@shitalbhandary